

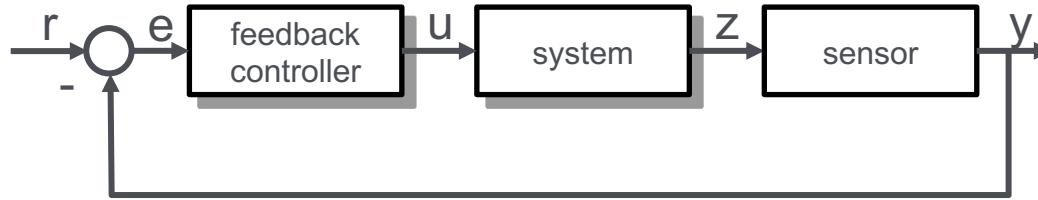
# Lecture 6: Nonlinear Least Squares

70003 Advanced Robotics

Dr Stefan Leutenegger

Teaching Assistants: Dr Masha Popovic,  
Sotiris Papatheodorou, Binbin Xu, and Nils Funk

# Proportional-Integral-Differential (PID) Control



$$u(t) = k_p e(t) + k_i \int_{t_0}^t e(\tau) d\tau + k_d \frac{de(t)}{dt}$$

**Proportional gain**  
reduces current error

**Integral gain**  
reduces long-term offset

**Differential gain**  
considers error trend,  
reduces overshoot

# Nonlinear Dynamic Systems

Typically, we will model as a **nonlinear continuous-time system** of the form:

$$\dot{\mathbf{x}}(t) = \mathbf{f}_c(\mathbf{x}(t), \mathbf{u}(t), \mathbf{w}(t)).$$

$$\mathbf{z}(t) = \mathbf{h}(\mathbf{x}(t)) + \mathbf{v}(t).$$

**Linearise:**  
around  $\bar{\mathbf{x}}, \bar{\mathbf{u}}$

The **linearised continuous-time system** can be handy for analysis:

$$\delta\dot{\mathbf{x}}(t) = \mathbf{F}_c\delta\mathbf{x}(t) + \mathbf{G}_c\delta\mathbf{u}(t) + \mathbf{L}_c\mathbf{w}(t).$$

$$\delta\mathbf{z}(t) = \mathbf{H}\delta\mathbf{x}(t) + \mathbf{v}(t).$$

**Discretise:** integrate  
from  $t_{k-1}$  to  $t_k$ .

$$\mathbf{x}_k = \mathbf{f}(\mathbf{x}_{k-1}, \mathbf{u}_k, \mathbf{w}_k).$$

$$\mathbf{z}_k = \mathbf{h}(\mathbf{x}_k) + \mathbf{v}_k.$$

**Linearise:**  
around  $\bar{\mathbf{x}}, \bar{\mathbf{u}}$

$$\delta\mathbf{x}_k = \mathbf{F}_k\delta\mathbf{x}_{k-1} + \mathbf{G}_k\delta\mathbf{u}_k + \mathbf{L}_k\mathbf{w}_k.$$

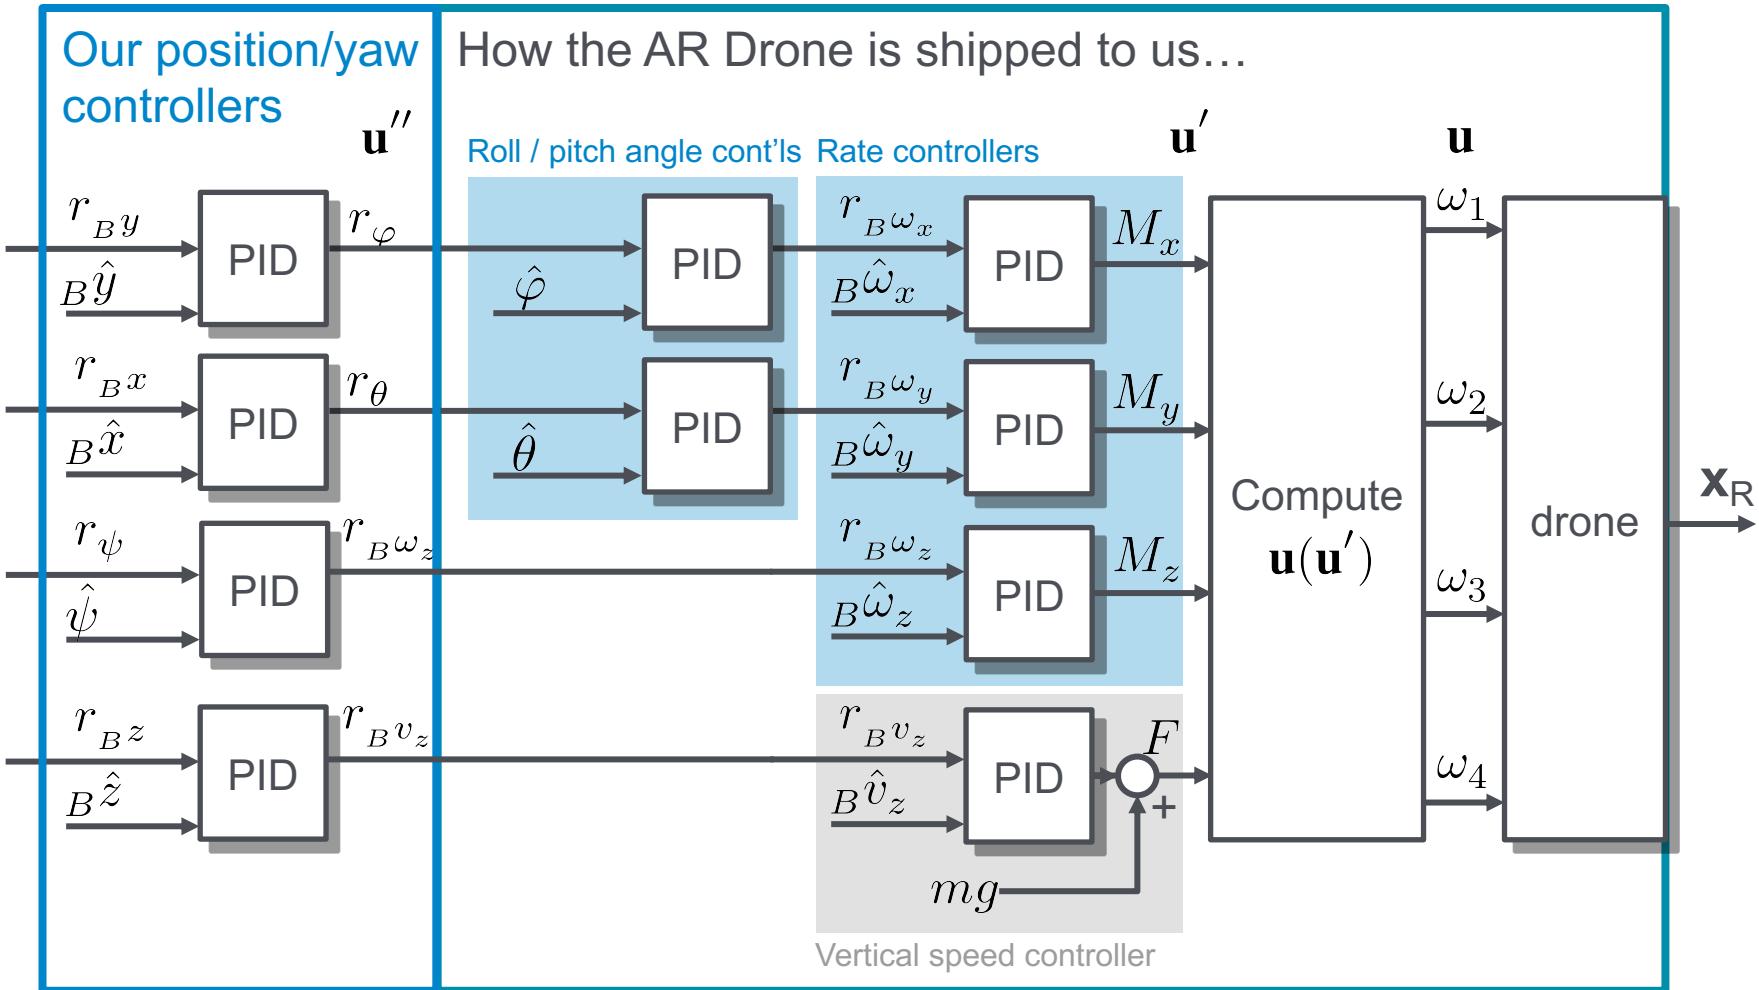
$$\delta\mathbf{z}_k = \mathbf{H}_k\delta\mathbf{x}_k + \mathbf{v}_k.$$

The **discrete-time nonlinear system** is used in our estimator implementations

Also the **linearised discrete-time system** is needed in the implementation

# PID Drone Control: Cascaded and Parallel SISO Loops

"Hats": state estimates



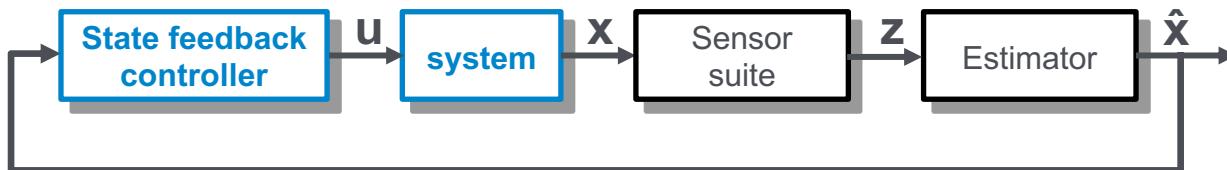
# Model-Based Control: Linear Quadratic Regulator (LQR)

Let's assume linear continuous-time dynamics

$$\dot{\mathbf{x}}(t) = \mathbf{F}_c \mathbf{x}(t) + \mathbf{G}_c \mathbf{u}(t),$$

That we would like to regulate to  $\mathbf{x}(t) = \mathbf{0}$ .

Now, let's try and find a **State-Feedback Control Law** of the form  
 $\mathbf{u}(t) = -\mathbf{K}\mathbf{x}(t)$ .



Now we want to minimise a cost functional

$$J = \int_t^\infty \underbrace{\mathbf{x}(\tau)^T \mathbf{Q} \mathbf{x}(\tau)}_{\text{Penalise deviation of state } \mathbf{x}(\tau) \text{ from } \mathbf{0}} + \underbrace{\mathbf{u}(\tau)^T \mathbf{R} \mathbf{u}(\tau)}_{\text{Penalise deviation of control input } \mathbf{u}(\tau) \text{ from } \mathbf{0}} d\tau.$$

(i.e. infinite-horizon;  
finite-horizon variants exist)

Penalise deviation  
of state  $\mathbf{x}(t)$  from  $\mathbf{0}$

Penalise deviation of  
control input  $\mathbf{u}(t)$  from  $\mathbf{0}$

## Model-Based Control: Linear Quadratic Regulator (LQR)

The solution (proof omitted here) is given by

$$\mathbf{K} = \mathbf{R}^{-1}(\mathbf{G}_c^T \mathbf{P}),$$

where  $\mathbf{P}$  is found by solving the so-called Algebraic Riccati Equation

$$\mathbf{F}_c^T \mathbf{P} + \mathbf{P} \mathbf{F}_c - (\mathbf{P} \mathbf{G}_c) \mathbf{R}^{-1}(\mathbf{G}_c^T \mathbf{P}) + \mathbf{Q} = \mathbf{0}.$$

This is a matrix equation that we cannot solve by hand normally.

But it's one line in Matlab: `K = lqr(F_c, G_c, Q, R);`

### Notes:

- Solving the (say with Matlab) for the feedback matrix  $\mathbf{K}$  is quite expensive, but is done **offline**! At runtime we only evaluate  $\mathbf{u}(t) = -\mathbf{K}\mathbf{x}(t)$ .
- We still have to set (i.e. **tune**) the weighting matrices  $\mathbf{Q}$  and  $\mathbf{R}$ . It's usually a good idea to use diagonal matrices.
- The optimality is nice, but **assumes no limits** on the system input  $\mathbf{u}$ . In practice, there will be limits and too high values on  $\mathbf{Q}$  will result in a bang-bang-style controller that is no longer optimal (or not even stable).
- We can use a **non-zero reference** point as well:  $\mathbf{u} = \mathbf{K}(\mathbf{r} - \mathbf{x})$ .

## LQR With Nonlinear Systems

We can do what we always do: use the linear approximation:

$$\delta \dot{\mathbf{x}}(t) = \mathbf{F}_c \delta \mathbf{x}(t) + \mathbf{G}_c \delta \mathbf{u}(t),$$

Around the current state (estimate)  $\bar{\mathbf{x}}$  and the current control input  $\bar{\mathbf{u}}$ .

**Problem:**  $\mathbf{F}_c$  and  $\mathbf{G}_c$  will be changing during runtime, but we don't want to solve the Matrix-Riccati equation in the control loop (too expensive)!

What we can do is **pre-compute** the feedback matrix  $\mathbf{K}$  for many linearisation points  $\bar{\mathbf{x}}$  and  $\bar{\mathbf{u}}$  as a **lookup-table**.

**Note:** all guarantees/optimality is lost here...

(But in practice this might still work very well)

# Probabilistic Estimation

We use Bayesian Statistics:

Measurements  $\mathbf{z}$  are modeled as samples from a **distribution**  $p(\mathbf{z}|\mathbf{x})$  given the variables  $\mathbf{x}$  (here: robot states plus possibly the map).

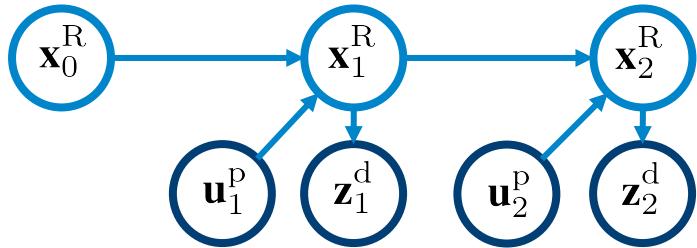
- **Maximum Likelihood (ML) Estimation:**  
What values for variables best explain all the measurements?
- **Maximum a Posteriori (MAP) Estimation:**  
What values for variables best explain all the measurements and a prior belief  $p(\mathbf{x})$  about those variables?

The EKF is one way of solving this problem.

Today we will see another one.

# Recursive Estimation and the Bayes Filter

- So far we have looked at batch estimation, i.e. estimating every variable at once.
- Recall, however, the typical structure of a robot state estimation problem:



$\mathbf{x}_k^R$  : the robot state at time step k.  
 $\mathbf{z}_k^d$  : a distance measurement.  
 $\mathbf{u}_k^p$  : wheel odometry (measurement).

$$p(\mathbf{x}_0^R, \dots, \mathbf{x}_N^R, \mathbf{z}_1^d, \dots, \mathbf{z}_N^d, \mathbf{u}_1^p, \dots, \mathbf{u}_N^p) = p(\mathbf{x}_0^R)p(\mathbf{x}_1^R|\mathbf{u}_1^p, \mathbf{x}_0^R)p(\mathbf{u}_1^p)p(\mathbf{z}_1^d|\mathbf{x}_1^R) \dots p(\mathbf{x}_N^R|\mathbf{u}_N^p, \mathbf{x}_{N-1}^R)p(\mathbf{u}_N^p)p(\mathbf{z}_N^d|\mathbf{x}_N^R).$$

Often, we are only interested in the **current state** (e.g. to control the robot) and not the whole (growing!) history of robot states (trajectory).

This can be done by alternating **prediction** (using a temporal model) and **update**:

## Bayes Filter – recursive MAP estimation:

- Initialise the prior distribution  $p(\mathbf{x}_0^R)$ .
- At every  $k^{\text{th}}$  iteration do the following two steps:

**1. Predict:**  $p(\mathbf{x}_k^R|\mathbf{u}_{1:k}^p, \mathbf{z}_{1:k-1}^d) = \int p(\mathbf{x}_k^R|\mathbf{u}_k^p, \mathbf{x}_{k-1}^R)p(\mathbf{x}_{k-1}^R|\mathbf{u}_{1:k-1}^p, \mathbf{z}_{1:k-1}^d)d\mathbf{x}_{k-1}^R$  **Prod./sum rule**

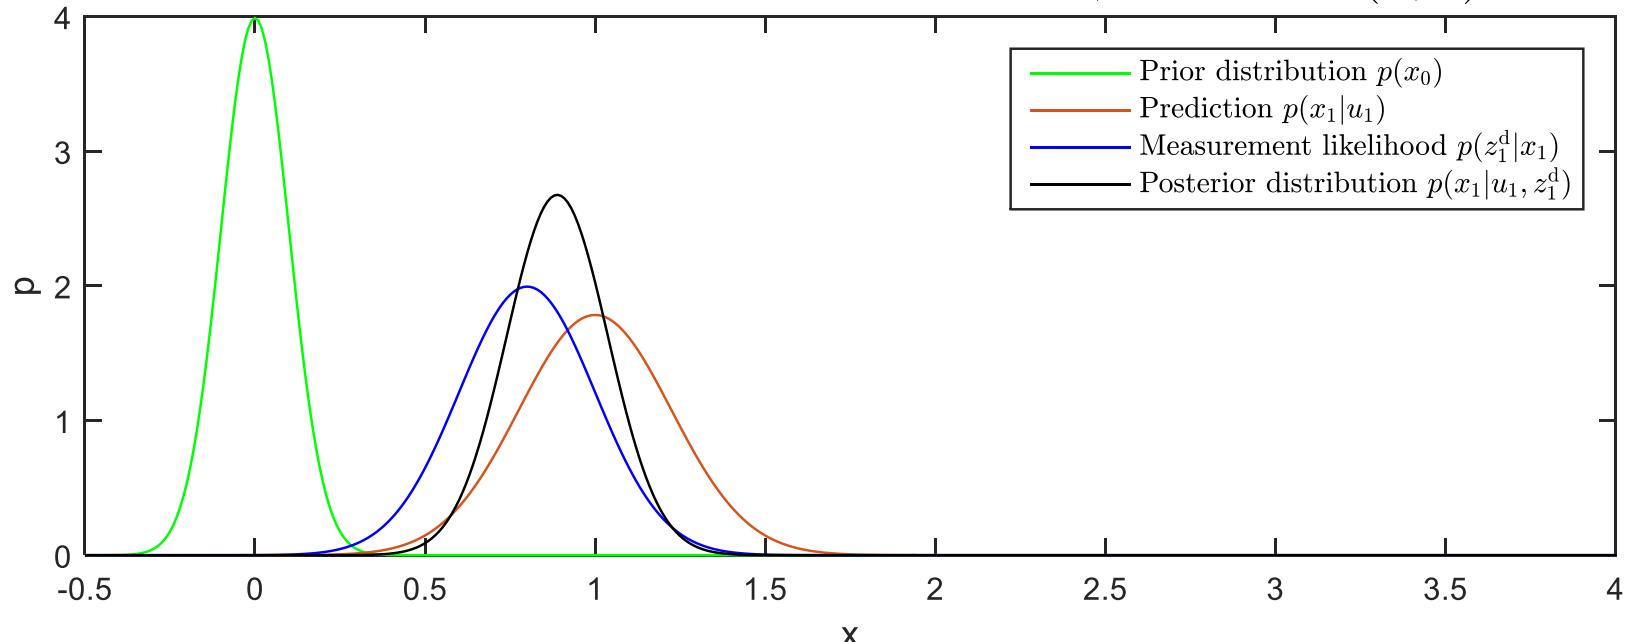
**2. Update:**  $p(\mathbf{x}_k^R|\mathbf{u}_{1:k}^p, \mathbf{z}_{1:k}^d) = \eta p(\mathbf{x}_k^R|\mathbf{u}_{1:k}^p, \mathbf{z}_{1:k-1}^d)p(\mathbf{z}_k^d|\mathbf{x}_k^R)$  **Bayes' theorem**

# 1D Kalman Filter (KF) Example

**KF: Bayes Filter assuming Gaussian distributions for all random variables as well as linear measurement and state transition models!**

**Example:** robot moving along the x-axis...

- State transition model:  $x_k = x_{k-1} + u + w$ ,  $w \sim \mathcal{N}(0, q)$  , where  $u^p$  is an odometry reading.
- Measurement updates (absolute position):  $z^d = x$ ,  $z^d \sim \mathcal{N}(0, r)$ .

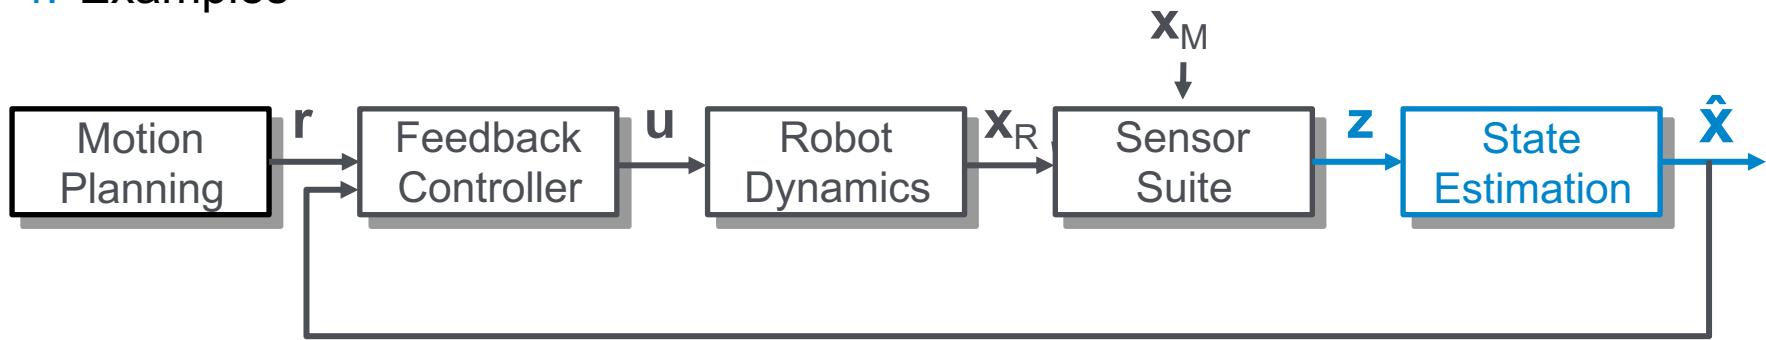


# Schedule: Lectures

What	Date	Topic
Lecture 1	18/01/21	Introduction, Problem Formulation, and Examples
Lecture 2	25/01/21	Representations and Sensors
Lecture 3	01/02/21	Kinematics and Temporal Models
Lecture 4	08/02/21	The Extended Kalman Filter
<b>Lecture 5</b>	<b>15/02/21</b>	<b>Feedback Control</b>
Lecture 6	22/02/21	Nonlinear Least Squares
Lecture 7	01/03/21	Vision-Based Simultaneous Localisation and Mapping
Lecture 8	08/03/21	Revision, Q&A

# Today

1. Linear Least Squares
2. Probabilistic Non-Linear Least Squares Estimation Problem
3. Optimising Non-Linear Least Squares Problems
4. Examples



# Linear Least Squares – Linear Function Fit

**Given:**

- $N$  data points  $\tilde{z}_i$  at  $t_i$ .
- Assumption that the data can be explained with a model  $z = (a_1 t + a_0)$ .

**Find:**

$\mathbf{x}^* = [a_0^*, a_1^*]^T$  minimising the sum of square differences

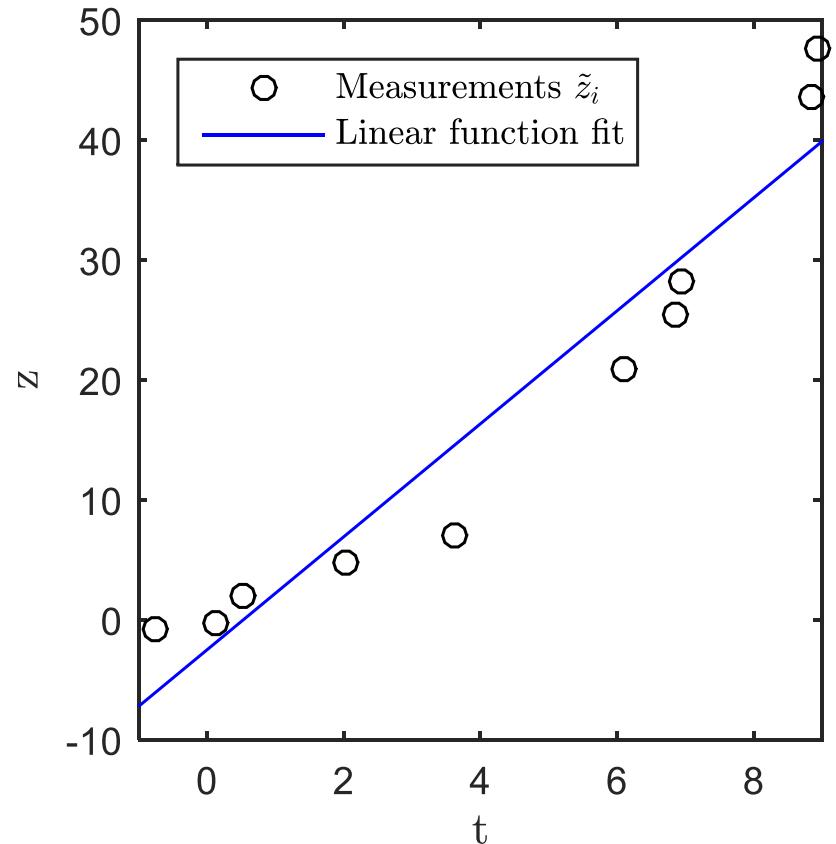
$$c(\mathbf{x}) = \frac{1}{2} \sum_{i=1}^N \|\tilde{z}_i - (a_1 t_i + a_0)\|_2^2.$$

**Solution:**

Re-write as  $c(\mathbf{x}) = \frac{1}{2} (\tilde{\mathbf{z}} - \mathbf{Hx})^T (\tilde{\mathbf{z}} - \mathbf{Hx})$ ,

$$\tilde{\mathbf{z}} = \begin{bmatrix} \tilde{z}_1 \\ \vdots \\ \tilde{z}_N \end{bmatrix}, \quad \mathbf{H} = \begin{bmatrix} 1 & t_1 \\ \vdots & \vdots \\ 1 & t_N \end{bmatrix}.$$

Necessary condition for optimality:



# Linear Least Squares – Polynomial Fit

**Given:**

- $N$  data points  $\tilde{z}_i$  at  $t_i$ .
- Assumption that the data can be explained with a model  $z = (a_2 t^2 + a_1 t + a_0)$ .

**Find:**

$\mathbf{x}^* = [a_0^*, a_1^*, a_2^*]^T$  minimising the sum of square differences

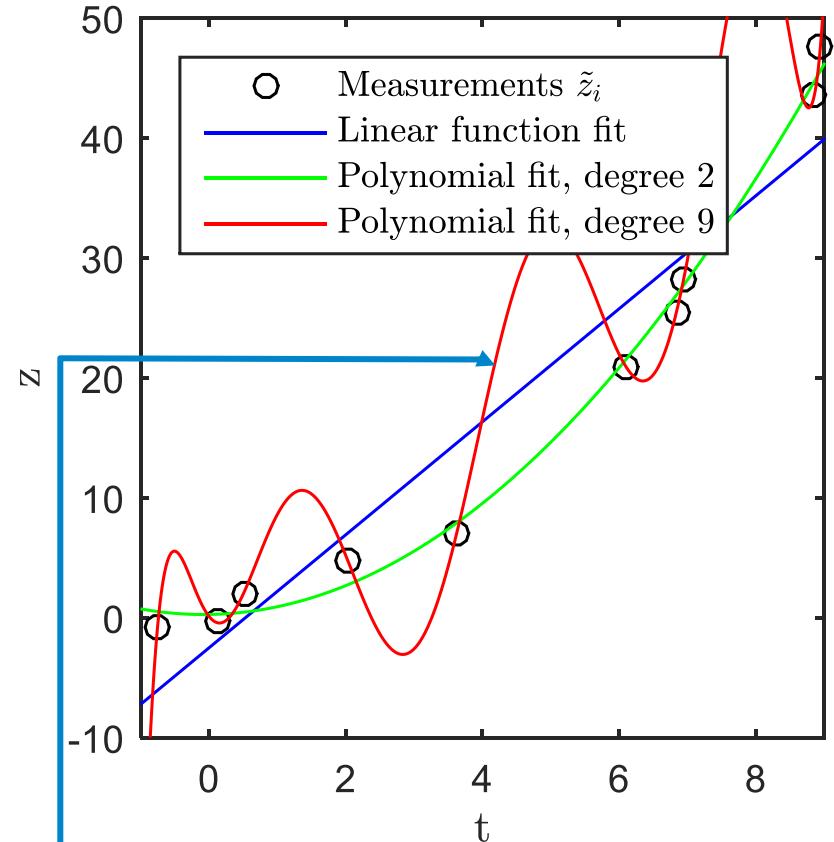
$$c(\mathbf{x}) = \frac{1}{2} \sum_{i=1}^N \|\tilde{z}_i - (a_2 t_i^2 + a_1 t_i + a_0)\|_2^2.$$

**Solution:**

Analogous to linear function fit, but with

$$\mathbf{H} = \begin{bmatrix} 1 & t_1 & t_1^2 \\ \vdots & \vdots & \vdots \\ 1 & t_N & t_N^2 \end{bmatrix}.$$

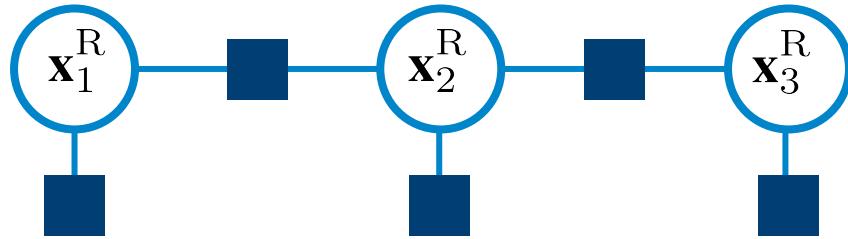
**Note:** This is a nonlinear function to fit, but still a **linear** regression problem!



**Overfitting:** the problem of explaining too little (noisy) data with too many parameters...

# Factor Graph – Differential Drive Robot Example

- All states of the trajectory are being estimated



■ Measurements: error terms

Measurements are of the form  $p(\mathbf{z}_i | \mathbf{x}_j)$  .

This will allow us to formulate measurement errors  $\mathbf{e}_i = \tilde{\mathbf{z}}_i - \mathbf{f}(\mathbf{x}_j)$ .

# The Probabilistic View on Estimation

**Given:**

- Variables to be estimated  $\mathbf{x}$ .
- $N$  independent measurements  $\tilde{\mathbf{z}}_i$  and associated measurement likelihood functions  $p(\mathbf{z}_i|\mathbf{x})$ .
- Optionally: a prior distribution for  $\mathbf{x}$ ,  $p(\mathbf{x})$ .

**Find:**

$$\mathbf{x}^* = \underset{N}{\operatorname{argmax}} p(\mathbf{x}|\mathbf{z}) \text{ with}$$

$$p(\mathbf{x}|\mathbf{z}) \propto p(\mathbf{z}|\mathbf{x}) = \prod_{i=1}^N p(\mathbf{z}_i|\mathbf{x}) \quad : \text{Maximum Likelihood, ML}$$

$$p(\mathbf{x}|\mathbf{z}) \propto p(\mathbf{z}|\mathbf{x})p(\mathbf{x}) = p(\mathbf{x}) \prod_{i=1}^N p(\mathbf{z}_i|\mathbf{x}) \quad : \text{Maximum a Posteriori, MAP}$$

# The Probabilistic View on Least Squares

**Now assume:**

- Gaussian measurement likelihood  $p(\mathbf{z}_i|\mathbf{x}) = \mathcal{N}(\mathbf{h}_i(\mathbf{x}), \mathbf{R}_i)$ .
- Optionally: Gaussian prior  $p(\mathbf{x}) = \mathcal{N}(\mathbf{x}_0, \mathbf{P})$ .

**Find:**  $\mathbf{x}^* = \operatorname{argmax} p(\mathbf{x}|\mathbf{z}) = \operatorname{argmin}(-\log(p(\mathbf{x}|\mathbf{z})))$

**ML Estimator:** minimise the weighted sum of least squares

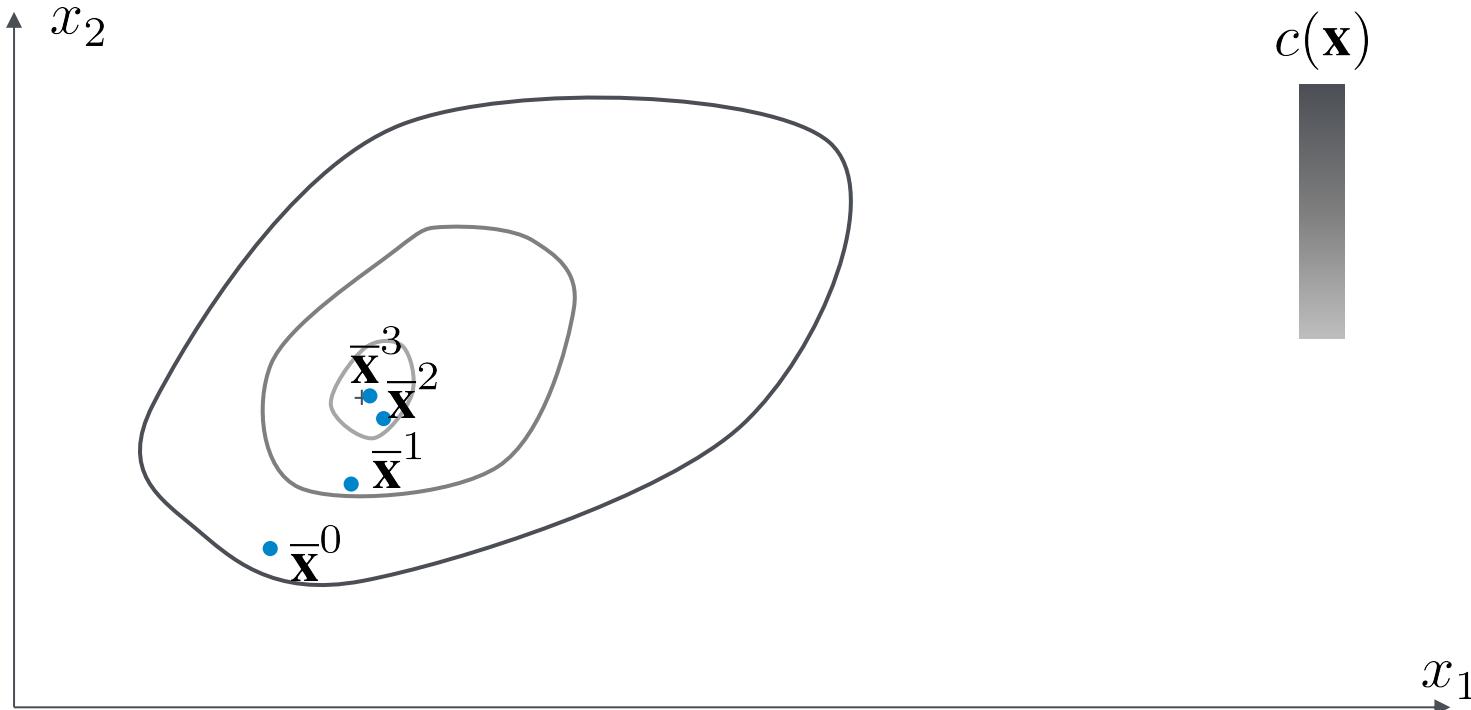
**MAP Estimator:** minimise the sum of least squares plus prior cost

# How To Solve Nonlinear Least Squares Problems [8,9]

**Linear case**  $\mathbf{h}_i(\mathbf{x}) = \mathbf{H}_i \mathbf{x}$

The solution is directly obtained using the normal equations.

**Nonlinear case: iterative methods starting at an initial guess for  $\mathbf{x}$**



# The Gauss-Newton Method

**Given:**

- Nonlinear error terms of the form  $e_i(\mathbf{x}) := \tilde{\mathbf{z}}_i - \mathbf{h}(\mathbf{x}) \sim \mathcal{N}(\mathbf{0}, \mathbf{R}_i)^*$ .

**Iterative solution to the associated nonlinear least squares problem:**

- Pick a suitable **starting point**  $\bar{\mathbf{x}}^0$ .
- Approximate the cost employing linear error terms  $e_i(\mathbf{x}) \approx e_i(\bar{\mathbf{x}}^k) + \mathbf{E}_i(\bar{\mathbf{x}}^k)\Delta\mathbf{x}$ , with

$$\mathbf{E}_i(\bar{\mathbf{x}}^k) = \frac{\partial}{\partial \mathbf{x}} e_i(\mathbf{x})|_{\bar{\mathbf{x}}^k} \quad (\text{note: typically sparse})$$

and formulate the **normal equations**:

$$\underbrace{\sum_i \mathbf{E}_i(\bar{\mathbf{x}}^k)^T \mathbf{R}_i^{-1} \mathbf{E}_i(\bar{\mathbf{x}}^k)}_{=:A} \Delta\mathbf{x} = -\underbrace{\sum_i \mathbf{E}_i(\bar{\mathbf{x}}^k)^T \mathbf{R}_i^{-1} e_i(\bar{\mathbf{x}}^k)}_{=:b}.$$

- Solve the normal equations for  $\Delta\mathbf{x}$  using e.g.

**Cholesky Decomposition**

(more on this step see Lecture 8).

- Update the current estimate as  $\bar{\mathbf{x}}^{k+1} = \bar{\mathbf{x}}^k + \Delta\mathbf{x}$ .
- Check **convergence**: if not OK, go to step 2.

# Levenberg-Marquardt

A trust-region method that interpolates between gradient-descent and Gauss-Newton.

1. Pick a **starting point**  $\bar{\mathbf{x}}^0$ , initial parameter  $\lambda^0 = \max \text{diag}(\mathbf{A})$  and  $\nu$ , e.g.  $\nu = 2$ .
2. Build the **modified** Gauss Newton system:  
$$(\mathbf{A} + \lambda \text{diag}(\mathbf{A}))\Delta\mathbf{x} = \mathbf{b}.$$
3. **Solve** it for  $\Delta\mathbf{x}$ .
4. **Update**
  - If cost is reduced:  $\bar{\mathbf{x}}^{k+1} = \bar{\mathbf{x}}^k + \Delta\mathbf{x}$  and expand trust region  $\lambda^{k+1} = \lambda^{k+1}/\nu$ .
  - Else: no update on  $\bar{\mathbf{x}}$ , contract the trust region  $\lambda^{k+1} = \lambda^{k+1}\nu$ ; go to Step 3
5. Check **convergence**: if not OK, go to step 2.

# Optimisation On Manifolds

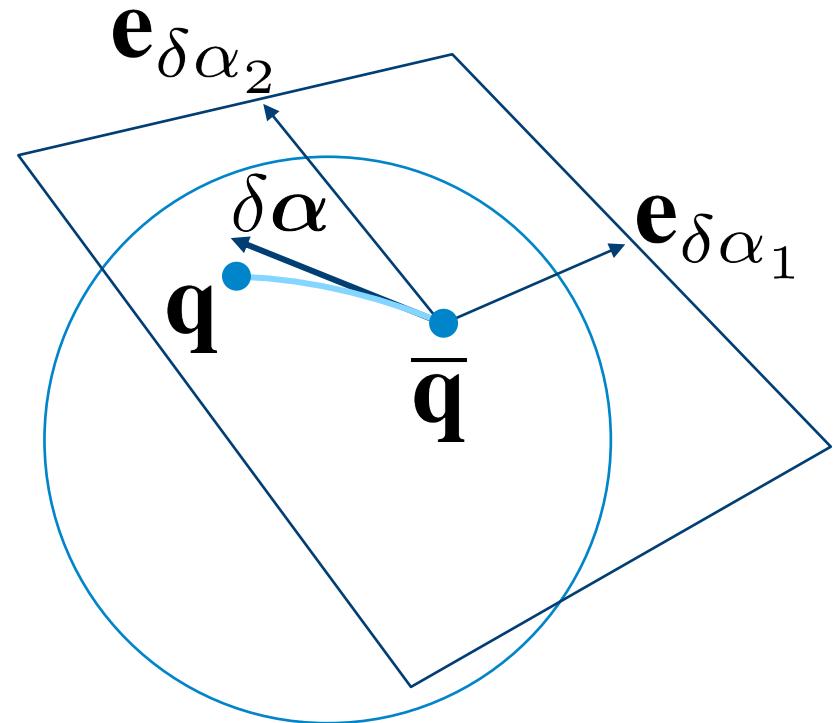
We may have to perform optimisations on manifolds...

**Example:** quaternions:

$$\delta \mathbf{q}(\delta \alpha) = \begin{bmatrix} \text{sinc} \left\| \frac{\delta \alpha}{2} \right\| \frac{\delta \alpha}{2} \\ \cos \left\| \frac{\delta \alpha}{2} \right\| \end{bmatrix}.$$

**Modifications** to solving optimisation problems:

- When computing Jacobians, compute them w.r.t.  $\delta \alpha \in \mathbb{R}^3$ .
- Compute the update  $\Delta \alpha$  and apply it as  $\bar{\mathbf{q}}^{k+1} = \delta \mathbf{q}(\Delta \alpha) \otimes \bar{\mathbf{q}}^k$ .



# Example: Camera Calibration

**Given:**

- A calibration pattern with  $j = 1, \dots, M$  observable points  ${}_T\mathbf{r}_{P_j}$  defined in the target frame  $\mathcal{F}_T$ .
- A set of images  $N$  of the calibration pattern taken by the same camera from different points of view. Assume the calibration points are identified in the image as  $\tilde{\mathbf{z}}_{i,P_j}$  (known 3D-2D correspondences).
- The pinhole camera projection model (with radial-tangential distortion)  
 $\mathbf{u} = \mathbf{k}(\mathbf{d}(\mathbf{p}({}_C\mathbf{r}_P)))$ .

**Find:**

- The most likely camera poses (extrinsics)  ${}_T\mathbf{r}_{C_i}, \mathbf{q}_{TC_i}$  per image.
- The most likely camera intrinsics parameters  $\mathbf{X}_c$ .

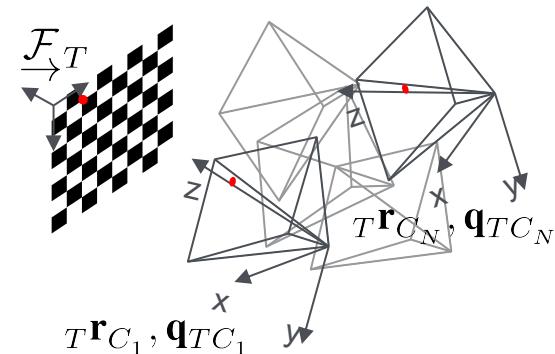
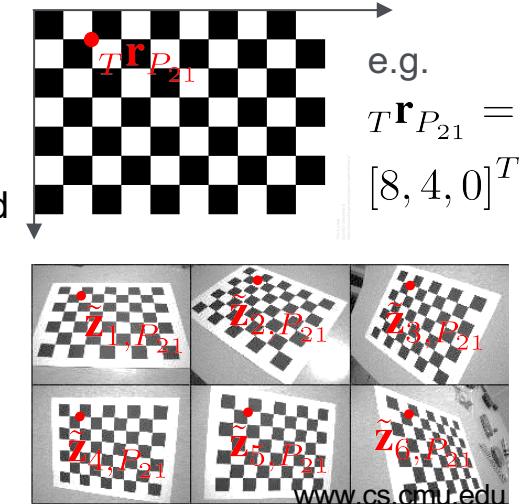
**Solution:**

We formulate the reprojection error per image and point:

$$\mathbf{e}_{i,j} = \tilde{\mathbf{z}}_{i,P_j} - \underbrace{\mathbf{k}(\mathbf{d}(\mathbf{p}(\mathbf{C}_{C_i T} {}_T\mathbf{r}_{P_j} - \mathbf{C}_{C_i T} {}_T\mathbf{r}_{C_i}))))}_{\text{Cam. projection}} \underbrace{- \tilde{\mathbf{z}}_{i,P_j}}_{\text{Transform point into camera frame}}$$

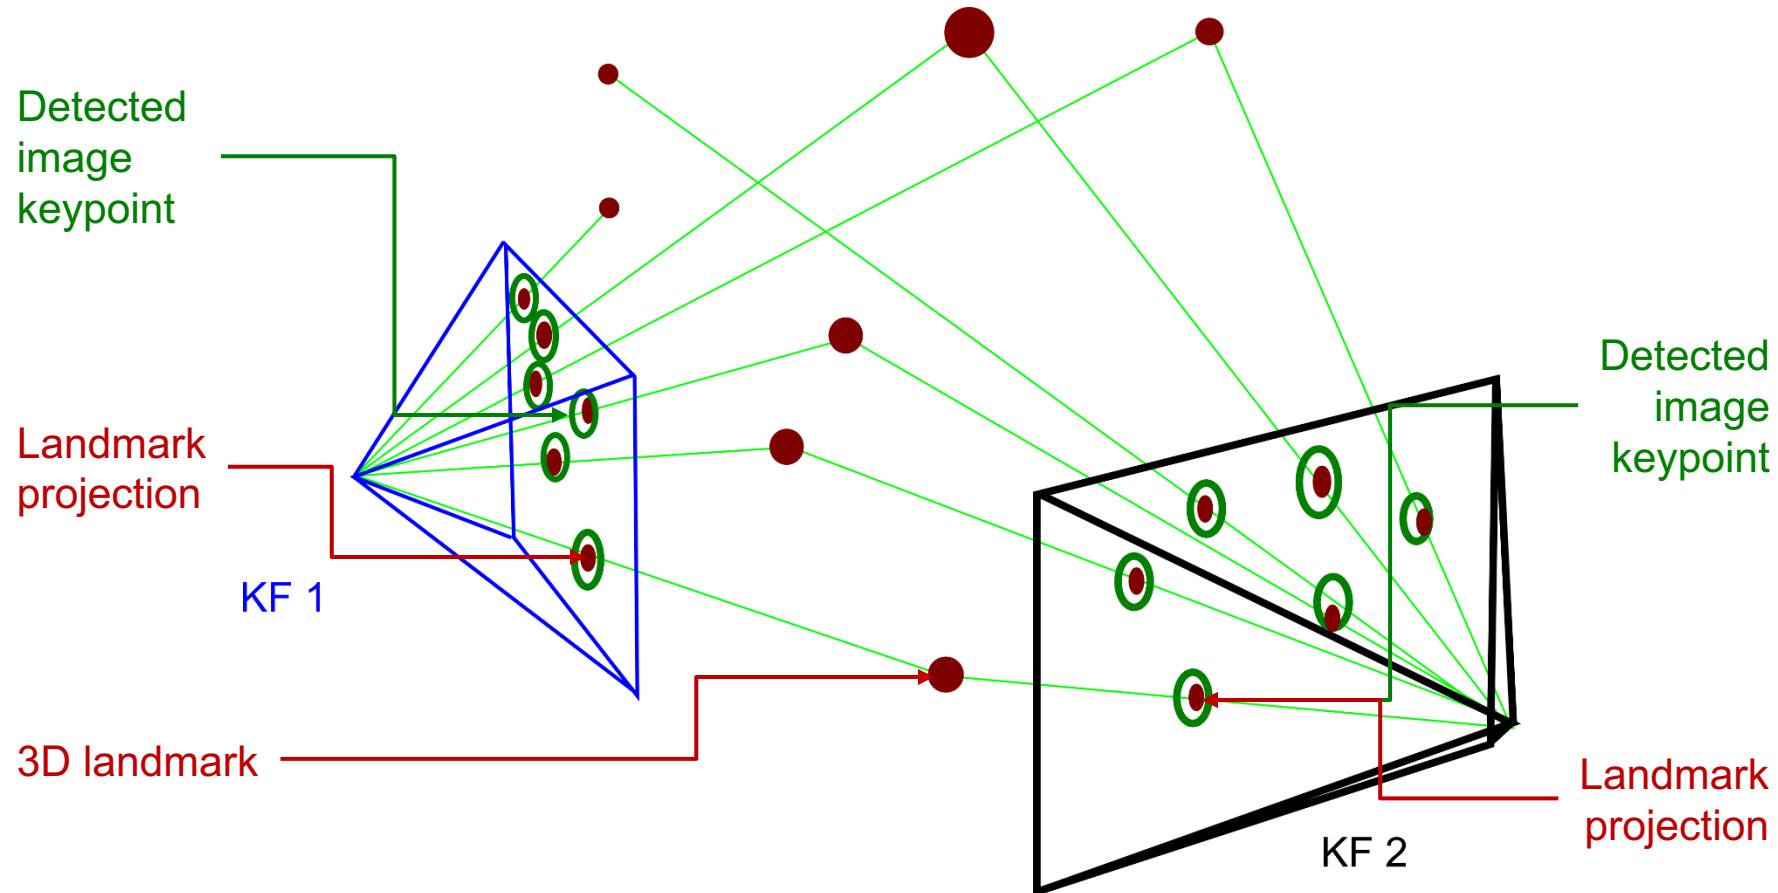
Minimise the overall reprojection cost  $c(\mathbf{X}_c, {}_T\mathbf{r}_{C_1}, \mathbf{q}_{TC_1}, \dots, {}_T\mathbf{r}_{C_N}, \mathbf{q}_{TC_N})$ , with  $\mathbf{R} = \text{diag}[\sigma^2, \sigma^2]$ , where  $\sigma$  is the detection uncertainty in pixels.

**Note:** we will need the Jacobians  $\frac{\partial \mathbf{e}_{i,j}}{\partial \mathbf{x}_c}$ ,  $\frac{\partial \mathbf{e}_{i,j}}{\partial {}_T\mathbf{r}_{C_i}}$  and  $\frac{\partial \mathbf{e}_{i,j}}{\partial \mathbf{q}_{TC_i}}$  (see lecture 4).



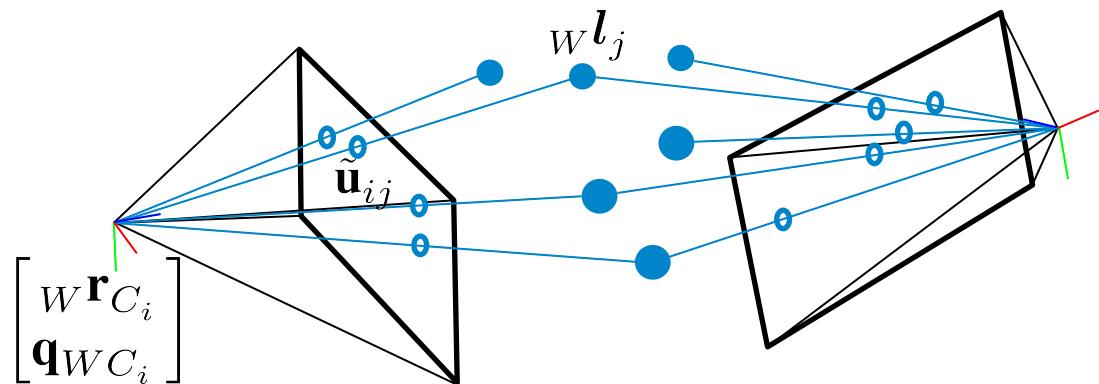
$$c(\mathbf{X}_c, {}_T\mathbf{r}_{C_1}, \mathbf{q}_{TC_1}, \dots, {}_T\mathbf{r}_{C_N}, \mathbf{q}_{TC_N}) = \sum_{i=1}^N \sum_{j=1}^M \mathbf{e}_{i,j}^T \mathbf{R}^{-1} \mathbf{e}_{i,j},$$

# The Visual SLAM Problem



## Bundle Adjustment: Visual SLAM as Optimisation Problem

**Given:** 2D detections  $\tilde{\mathbf{u}}_{ij}$  of corresponding unknown 3D landmarks  ${}_W\mathbf{l}_j$  in images  $i = 1 \dots N$ .



**Find:** 3D landmark positions  ${}_W\mathbf{l}_j$  and 6D camera poses  $[{}_W\mathbf{r}_{C_i}, \mathbf{q}_{WC_i}]^T$  where the images were taken.

**Solution:** minimise the reprojection error

$$\mathbf{e}_{ij}({}_W\mathbf{r}_{C_i}, \mathbf{q}_{WC_i}, {}_W\mathbf{l}_j) = \tilde{\mathbf{u}} - \mathbf{u}(T_{C_i W} {}_W\mathbf{l}_j),$$

over all images

$$\min \sum_{i=1}^N \sum_{j \in \mathcal{J}(i)} \|\mathbf{e}_{ij}\|^2,$$

Set of landmark indices as visible in camera  $i$ .

applying Levenberg-Marquardt.

## Jacobians needed for the Bundle Adjustment Problem

We will need **Jacobians** of  $\mathbf{e}_{ij}({}_W\mathbf{r}_{C_i}, \mathbf{q}_{WC_i}, {}_W\mathbf{l}_j) = \tilde{\mathbf{u}} - \mathbf{u}(T_{C_i W} {}_W\mathbf{l}_j)$ .

Let's rewrite first:  $\mathbf{e}_{ij} = \tilde{\mathbf{u}} - \mathbf{u}(T_{WC_i}^{-1} {}_W\mathbf{l}_j) = \tilde{\mathbf{u}} - \mathbf{u}(\mathbf{C}_{WC_i}^T {}_W\mathbf{l}_j - \mathbf{C}_{WC_i}^T {}_W\mathbf{r}_{C_i})$ .

We will employ the chain rule. So, we first need  $\frac{\partial}{\partial {}_{C_i}\mathbf{l}} \mathbf{u}({}_{C_i}\mathbf{l}) = \frac{\partial}{\partial {}_{C_i}\mathbf{l}} \mathbf{k}(\mathbf{d}(\mathbf{p}({}_{C_i}\mathbf{l})))$ .

$$\mathbf{U}({}_{C_i}\bar{\mathbf{l}}_j) := \frac{\partial}{\partial {}_{C_i}\bar{\mathbf{l}}_j} \mathbf{u}({}_{C_i}\mathbf{l}_j) = \begin{bmatrix} f_u & 0 \\ 0 & f_v \end{bmatrix} \mathbf{D} \begin{bmatrix} \frac{1}{\bar{l}_3} & 0 & -\frac{\bar{l}_1}{\bar{l}_3^2} \\ 0 & \frac{1}{\bar{l}_3} & -\frac{\bar{l}_2}{\bar{l}_3^2} \end{bmatrix}$$

Jacobian of distortion

Now:

$$\frac{\partial {}_{C_i}\mathbf{l}_j}{\partial {}_W\mathbf{r}_{C_i}} = -\bar{\mathbf{C}}_{WC_i}^T, \quad \frac{\partial {}_{C_i}\mathbf{l}_j}{\partial \delta\boldsymbol{\alpha}_i} = \bar{\mathbf{C}}_{WC_i}^T [{}_W\bar{\mathbf{l}}_j - {}_W\bar{\mathbf{r}}_{C_i}]^\times, \quad \frac{\partial {}_{C_i}\mathbf{l}_j}{\partial {}_W\bar{\mathbf{l}}_j} = \bar{\mathbf{C}}_{WC_i}^T.$$

Now we have everything to build the Gauss-Newton System of equations:

$$\underbrace{\sum_{ij} \mathbf{E}_{ij}(\bar{\mathbf{x}})^T \mathbf{R}_{ij}^{-1} \mathbf{E}_{ij}(\bar{\mathbf{x}})}_{=:A} \Delta \mathbf{x} = -\underbrace{\sum_{ij} \mathbf{E}_{ij}(\bar{\mathbf{x}})^T \mathbf{R}_{ij}^{-1} \mathbf{e}_{ij}(\bar{\mathbf{x}})}_{=:b},$$

$\mathbf{E}_{ij}(\bar{\mathbf{x}}) = \frac{\partial}{\partial \mathbf{x}} \mathbf{e}_{ij}(\mathbf{x}) \Big|_{\bar{\mathbf{x}}}.$

Jacobian of reproj. error w.r.t.  
all the variables (lots of zeros)

# Structure of the Bundle Adjustment Problem

The **Gauss-Newton system of equations** smartly (additive per reproj. error):

$$\begin{bmatrix} \frac{\partial \mathbf{e}_{ij}}{\partial_W \mathbf{r}_{C_i}}, \frac{\partial \mathbf{e}_{ij}}{\partial \delta \boldsymbol{\alpha}_i} \end{bmatrix}^T \mathbf{W}_{ij} \begin{bmatrix} \frac{\partial \mathbf{e}_{ij}}{\partial_W \mathbf{r}_{C_i}}, \frac{\partial \mathbf{e}_{ij}}{\partial \delta \boldsymbol{\alpha}_i} \end{bmatrix} = \begin{bmatrix} \frac{\partial \mathbf{e}_{ij}}{\partial_W \mathbf{r}_{C_i}}, \frac{\partial \mathbf{e}_{ij}}{\partial \delta \boldsymbol{\alpha}_i} \end{bmatrix}^T \mathbf{W}_{ij} \frac{\partial \mathbf{e}_{ij}}{\partial_W \mathbf{l}_j} - \begin{bmatrix} \frac{\partial \mathbf{e}_{ij}}{\partial_W \mathbf{r}_{C_i}}, \frac{\partial \mathbf{e}_{ij}}{\partial \delta \boldsymbol{\alpha}_i} \end{bmatrix}^T \mathbf{W}_{ij} \bar{\mathbf{e}}_{ij}$$

$\mathbf{A}$

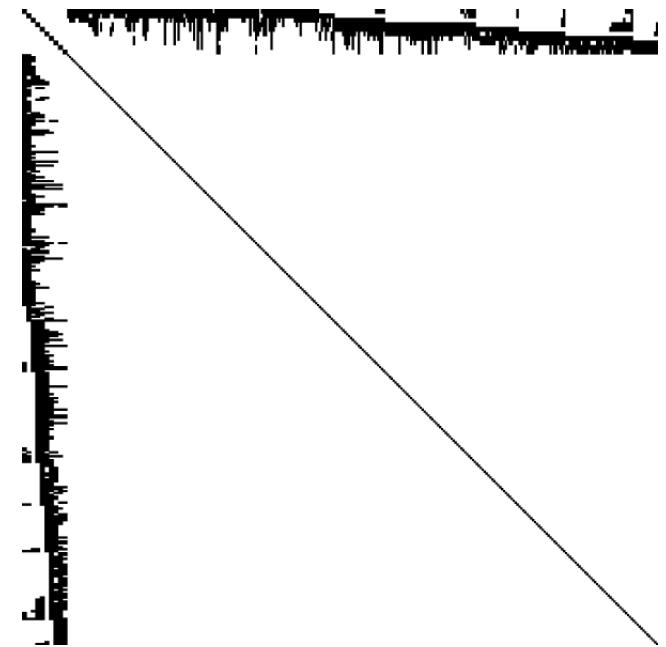
Weights:  
inverse  
detection  
covariance  
(can be  
identity here)

Error evaluated  
at linearisation  
point

# Some Notes on Bundle Adjustment

- Typical problems contain many more points than poses
- Due to the sparsity of the linear equation system, it needs to be solved in a with some tricks, e.g.
  - Use Schur-Complement elimination of landmarks
  - Use sparse Cholesky composition
- In practise, **use an off-the-shelf solver** like
  - ceres: <http://ceres-solver.org/>
  - g2o:  
<https://github.com/RainerKuemmerle/g2o>

This means, you will only have to implement the error function and Jacobians!



Example matrix A of a modestly sized problem.

[https://en.wikipedia.org/wiki/Bundle\\_adjustment](https://en.wikipedia.org/wiki/Bundle_adjustment)

# Batch Optimisation vs. Recursive Estimation (Filtering)

## Batch Optimisation

- Estimate all states ever at once.
- No need for a temporal model, but can employ one.
- Enables re-linearization, leading to highest possible accuracy.
- However: when estimating a time-varying state, this means inserting variables for each time-step (i.e. the problem will grow with time)!

## Recursive Estimation

- Estimate only the current state and maintain some distribution.
- Needs a temporal model for prediction, which is alternated with measurement updates.
- Constant state vector size, constant complexity.
- Re-linearization for old state is impossible (reduced accuracy).
- Delayed measurements cannot be included.

# Any Questions?

See you on Thursday at 1pm for the practical