

Lecture 7: Vision-Based Simultaneous Localisation and Mapping

70003 Advanced Robotics

Dr Stefan Leutenegger

Teaching Assistants: Dr Masha Popovic,
Sotiris Papatheodorou, Binbin Xu, and Nils Funk

Linear Least Squares – Polynomial Fit

Given:

- N data points \tilde{z}_i at t_i .
- Assumption that the data can be explained with a model $z = (a_2 t^2 + a_1 t + a_0)$.

Find:

$\mathbf{x}^* = [a_0^*, a_1^*, a_2^*]^T$ minimising the sum of square differences

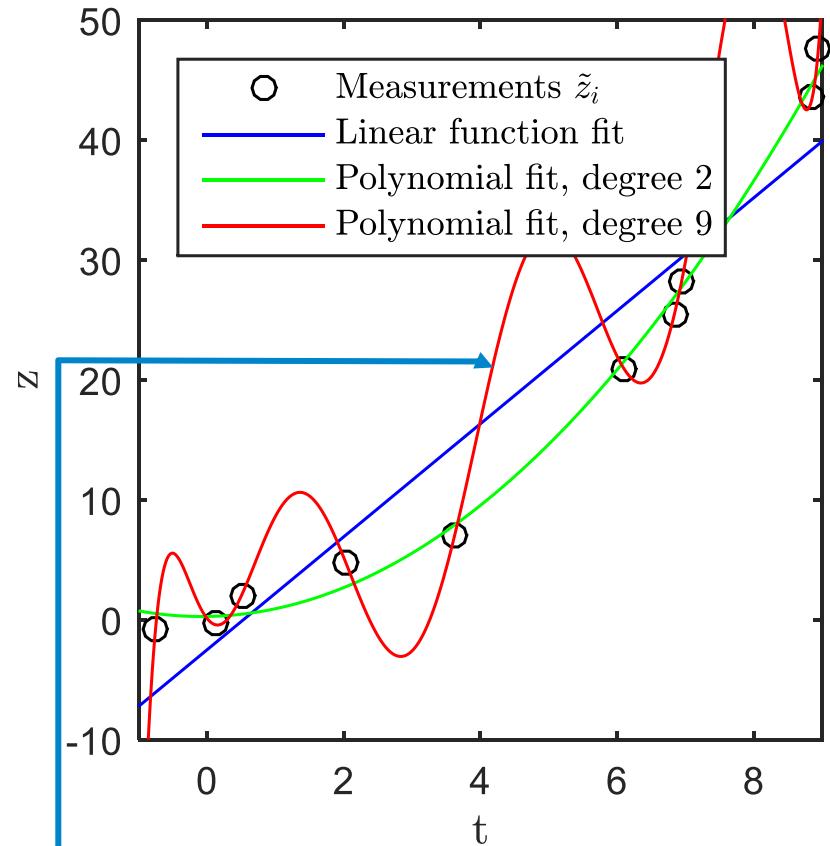
$$c(\mathbf{x}) = \frac{1}{2} \sum_{i=1}^N \|\tilde{z}_i - (a_2 t_i^2 + a_1 t_i + a_0)\|_2^2.$$

Solution:

Analogous to linear function fit, but with

$$\mathbf{H} = \begin{bmatrix} 1 & t_1 & t_1^2 \\ \vdots & \vdots & \vdots \\ 1 & t_N & t_N^2 \end{bmatrix}.$$

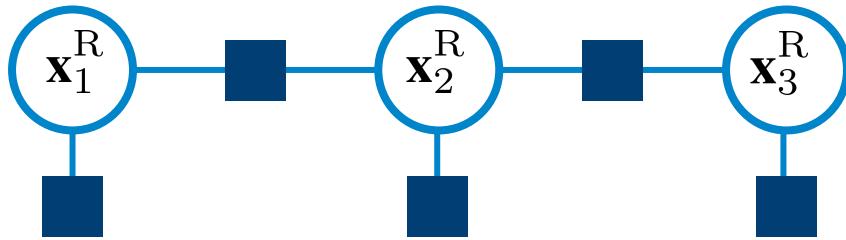
Note: This is a nonlinear function to fit, but still a **linear** regression problem!



Overfitting: the problem of explaining too little (noisy) data with too many parameters...

Factor Graph – Differential Drive Robot Example

- All states of the trajectory are being estimated



■ Measurements: error terms

Measurements are of the form $p(\mathbf{z}_i | \mathbf{x}_j)$.

This will allow us to formulate measurement errors $\mathbf{e}_i = \tilde{\mathbf{z}}_i - \mathbf{f}(\mathbf{x}_j)$.

The Probabilistic View on Least Squares

Now assume:

- Gaussian measurement likelihood $p(\mathbf{z}_i|\mathbf{x}) = \mathcal{N}(\mathbf{h}_i(\mathbf{x}), \mathbf{R}_i)$.
- Optionally: Gaussian prior $p(\mathbf{x}) = \mathcal{N}(\mathbf{x}_0, \mathbf{P})$.

Find: $\mathbf{x}^* = \operatorname{argmax} p(\mathbf{x}|\mathbf{z}) = \operatorname{argmin}(-\log(p(\mathbf{x}|\mathbf{z})))$

ML Estimator: minimise the weighted sum of least squares

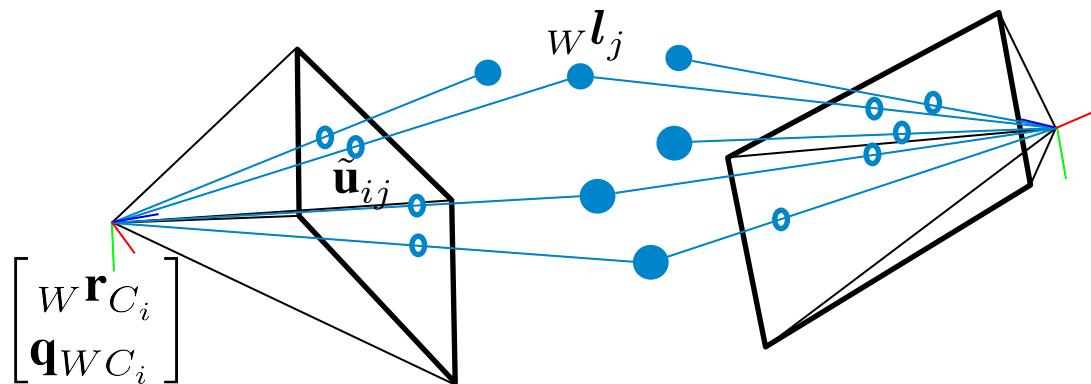
$$\begin{aligned}\mathbf{x}^* &= \operatorname{argmin}(-\log(p(\mathbf{x}|\mathbf{z}))) = \operatorname{argmin}(-\log(p(\mathbf{z}|\mathbf{x}))) \\ &= \operatorname{argmin} \left(-\log \left(\prod_{i=1}^N p(\mathbf{z}_i|\mathbf{x}) \right) \right) = \operatorname{argmin} \left(-\log \left(\prod_{i=1}^N \exp \left(-\frac{1}{2} (\tilde{\mathbf{z}}_i - \mathbf{h}_i(\mathbf{x}))^\top \mathbf{R}_i^{-1} (\tilde{\mathbf{z}}_i - \mathbf{h}_i(\mathbf{x})) \right) \right) \right) \\ &= \operatorname{argmin} \frac{1}{2} \sum_{i=1}^N (\tilde{\mathbf{z}}_i - \mathbf{h}_i(\mathbf{x}))^\top \mathbf{R}_i^{-1} (\tilde{\mathbf{z}}_i - \mathbf{h}_i(\mathbf{x})).\end{aligned}$$

MAP Estimator: minimise the sum of least squares plus prior cost

$$\begin{aligned}\mathbf{x}^* &= \operatorname{argmin}(-\log(p(\mathbf{x}|\mathbf{z}))) = \operatorname{argmin}(-\log(p(\mathbf{z}|\mathbf{x})p(\mathbf{x}))) \\ &= \operatorname{argmin} \left(\frac{1}{2} (\mathbf{x}_0 - \mathbf{x})^\top \mathbf{P}^{-1} (\mathbf{x}_0 - \mathbf{x}) + \frac{1}{2} \sum_{i=1}^N (\tilde{\mathbf{z}}_i - \mathbf{h}_i(\mathbf{x}))^\top \mathbf{R}_i^{-1} (\tilde{\mathbf{z}}_i - \mathbf{h}_i(\mathbf{x})) \right).\end{aligned}$$

Bundle Adjustment: Visual SLAM as Optimisation Problem

Given: 2D detections $\tilde{\mathbf{u}}_{ij}$ of corresponding unknown 3D landmarks ${}_W\mathbf{l}_j$ in images $i = 1 \dots N$.



Find: 3D landmark positions ${}_W\mathbf{l}_j$ and 6D camera poses $[{}_W\mathbf{r}_{C_i}, {}_W\mathbf{q}_{C_i}]^T$ where the images were taken.

Solution: minimise the reprojection error

$$\mathbf{e}_{ij}({}_W\mathbf{r}_{C_i}, {}_W\mathbf{q}_{C_i}, {}_W\mathbf{l}_j) = \tilde{\mathbf{u}} - \mathbf{u}({}_W\mathbf{T}_{C_i W} {}_W\mathbf{l}_j),$$

over all images

$$\min \sum_{i=1}^N \sum_{j \in \mathcal{J}(i)} \|\mathbf{e}_{ij}\|^2,$$

Set of landmark indices as visible in camera i.

applying Levenberg-Marquardt.

Jacobians needed for the Bundle Adjustment Problem

We will need **Jacobians** of $\mathbf{e}_{ij}({}_W\mathbf{r}_{C_i}, \mathbf{q}_{WC_i}, {}_W\mathbf{l}_j) = \tilde{\mathbf{u}} - \mathbf{u}(T_{C_i W} {}_W\mathbf{l}_j)$.

Let's rewrite first: $\mathbf{e}_{ij} = \tilde{\mathbf{u}} - \mathbf{u}(T_{WC_i}^{-1} {}_W\mathbf{l}_j) = \tilde{\mathbf{u}} - \mathbf{u}(\mathbf{C}_{WC_i}^T {}_W\mathbf{l}_j - \mathbf{C}_{WC_i}^T {}_W\mathbf{r}_{C_i})$.

We will employ the chain rule. So, we first need $\frac{\partial}{\partial {}_{C_i}\mathbf{l}} \mathbf{u}({}_{C_i}\mathbf{l}) = \frac{\partial}{\partial {}_{C_i}\mathbf{l}} \mathbf{k}(\mathbf{d}(\mathbf{p}({}_{C_i}\mathbf{l})))$.

$$\mathbf{U}({}_{C_i}\bar{\mathbf{l}}_j) := \frac{\partial}{\partial {}_{C_i}\bar{\mathbf{l}}_j} \mathbf{u}({}_{C_i}\mathbf{l}_j) = \begin{bmatrix} f_u & 0 \\ 0 & f_v \end{bmatrix} \mathbf{D} \begin{bmatrix} \frac{1}{\bar{l}_3} & 0 & -\frac{\bar{l}_1}{\bar{l}_3^2} \\ 0 & \frac{1}{\bar{l}_3} & -\frac{\bar{l}_2}{\bar{l}_3^2} \end{bmatrix}$$

Jacobian of distortion

Now:

$$\frac{\partial {}_{C_i}\mathbf{l}_j}{\partial {}_W\mathbf{r}_{C_i}} = -\bar{\mathbf{C}}_{WC_i}^T, \quad \frac{\partial {}_{C_i}\mathbf{l}_j}{\partial \delta\boldsymbol{\alpha}_i} = \bar{\mathbf{C}}_{WC_i}^T [{}_W\bar{\mathbf{l}}_j - {}_W\bar{\mathbf{r}}_{C_i}]^\times, \quad \frac{\partial {}_{C_i}\mathbf{l}_j}{\partial {}_W\bar{\mathbf{l}}_j} = \bar{\mathbf{C}}_{WC_i}^T.$$

Now we have everything to build the Gauss-Newton System of equations:

$$\underbrace{\sum_{ij} \mathbf{E}_{ij}(\bar{\mathbf{x}})^T \mathbf{R}_{ij}^{-1} \mathbf{E}_{ij}(\bar{\mathbf{x}})}_{=:A} \Delta \mathbf{x} = -\underbrace{\sum_{ij} \mathbf{E}_{ij}(\bar{\mathbf{x}})^T \mathbf{R}_{ij}^{-1} \mathbf{e}_{ij}(\bar{\mathbf{x}})}_{=:b},$$

$\mathbf{E}_{ij}(\bar{\mathbf{x}}) = \frac{\partial}{\partial \mathbf{x}} \mathbf{e}_{ij}(\mathbf{x}) \Big|_{\bar{\mathbf{x}}}.$

Jacobian of reproj. error w.r.t.
all the variables (lots of zeros)

Structure of the Bundle Adjustment Problem

The **Gauss-Newton system of equations** smartly (additive per reproj. error):

$$\begin{bmatrix} \frac{\partial \mathbf{e}_{ij}}{\partial_W \mathbf{r}_{C_i}}, \frac{\partial \mathbf{e}_{ij}}{\partial \delta \boldsymbol{\alpha}_i} \end{bmatrix}^T \mathbf{W}_{ij} \begin{bmatrix} \frac{\partial \mathbf{e}_{ij}}{\partial_W \mathbf{r}_{C_i}}, \frac{\partial \mathbf{e}_{ij}}{\partial \delta \boldsymbol{\alpha}_i} \end{bmatrix} = \begin{bmatrix} \frac{\partial \mathbf{e}_{ij}}{\partial_W \mathbf{r}_{C_i}}, \frac{\partial \mathbf{e}_{ij}}{\partial \delta \boldsymbol{\alpha}_i} \end{bmatrix}^T \mathbf{W}_{ij} \frac{\partial \mathbf{e}_{ij}}{\partial_W \mathbf{l}_j} - \begin{bmatrix} \frac{\partial \mathbf{e}_{ij}}{\partial_W \mathbf{r}_{C_i}}, \frac{\partial \mathbf{e}_{ij}}{\partial \delta \boldsymbol{\alpha}_i} \end{bmatrix}^T \mathbf{W}_{ij} \bar{\mathbf{e}}_{ij}$$

A

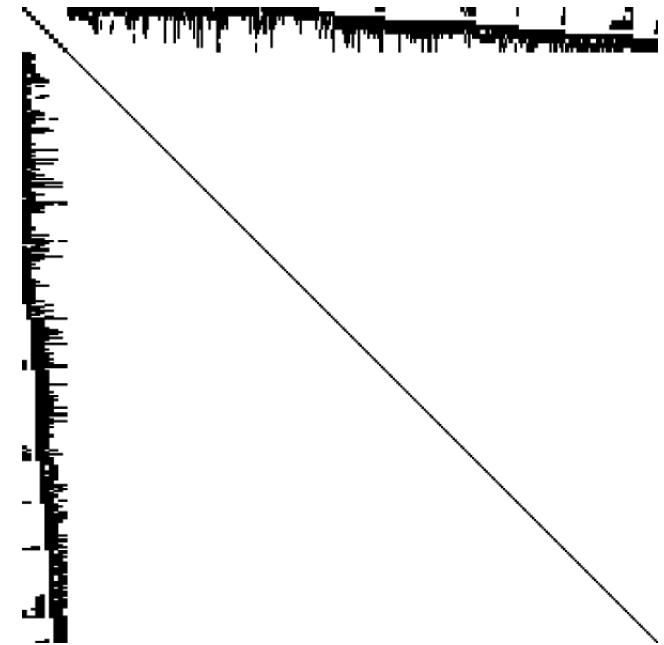
Weights:
inverse
detection
covariance
(can be
identity here)

Error evaluated
at linearisation
point

Some Notes on Bundle Adjustment

- Typical problems contain many more points than poses
- Due to the sparsity of the linear equation system, it needs to be solved in a with some tricks, e.g.
 - Use Schur-Complement elimination of landmarks
 - Use sparse Cholesky composition
- In practise, **use an off-the-shelf solver** like
 - ceres: <http://ceres-solver.org/>
 - g2o:
<https://github.com/RainerKuemmerle/g2o>

This means, you will only have to implement the error function and Jacobians!



Example matrix A of a modestly sized problem.

https://en.wikipedia.org/wiki/Bundle_adjustment

Batch Optimisation vs. Recursive Estimation (Filtering)

Batch Optimisation

- Estimate all states ever at once.
- No need for a temporal model, but can employ one.
- Enables re-linearization, leading to highest possible accuracy.
- However: when estimating a time-varying state, this means inserting variables for each time-step (i.e. the problem will grow with time)!

Recursive Estimation

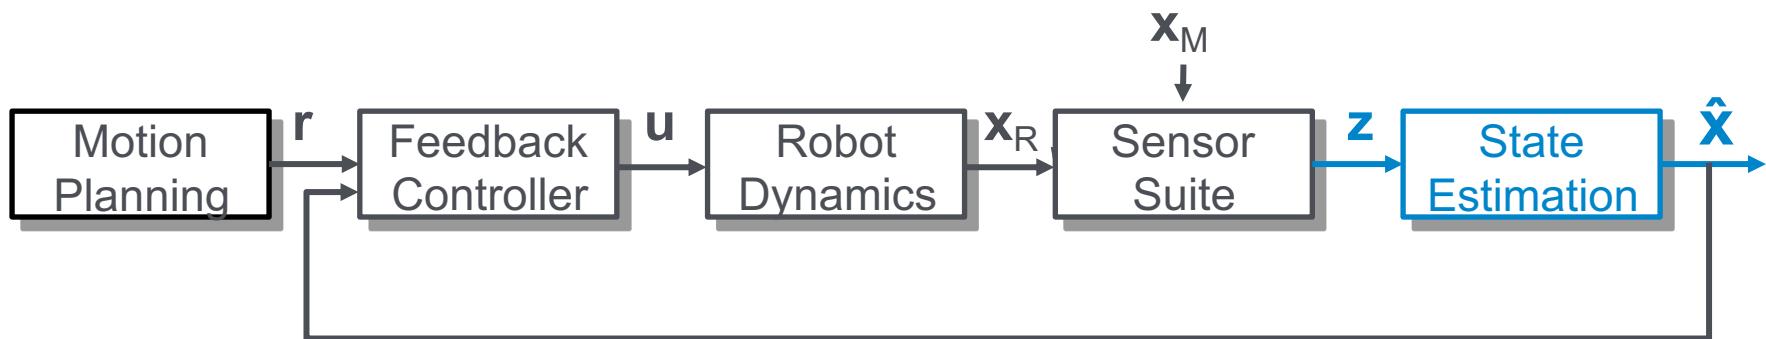
- Estimate only the current state and maintain some distribution.
- Needs a temporal model for prediction, which is alternated with measurement updates.
- Constant state vector size, constant complexity.
- Re-linearization for old state is impossible (reduced accuracy).
- Delayed measurements cannot be included.

Schedule: Lectures

What	Date	Topic
Lecture 1	18/01/21	Introduction, Problem Formulation, and Examples
Lecture 2	25/01/21	Representations and Sensors
Lecture 3	01/02/21	Kinematics and Temporal Models
Lecture 4	08/02/21	The Extended Kalman Filter
Lecture 5	15/02/21	Feedback Control
Lecture 6	22/02/21	Nonlinear Least Squares
Lecture 7	01/03/21	Vision-based Simultaneous Localisation and Mapping
Lecture 8	08/03/21	Revision, Q&A

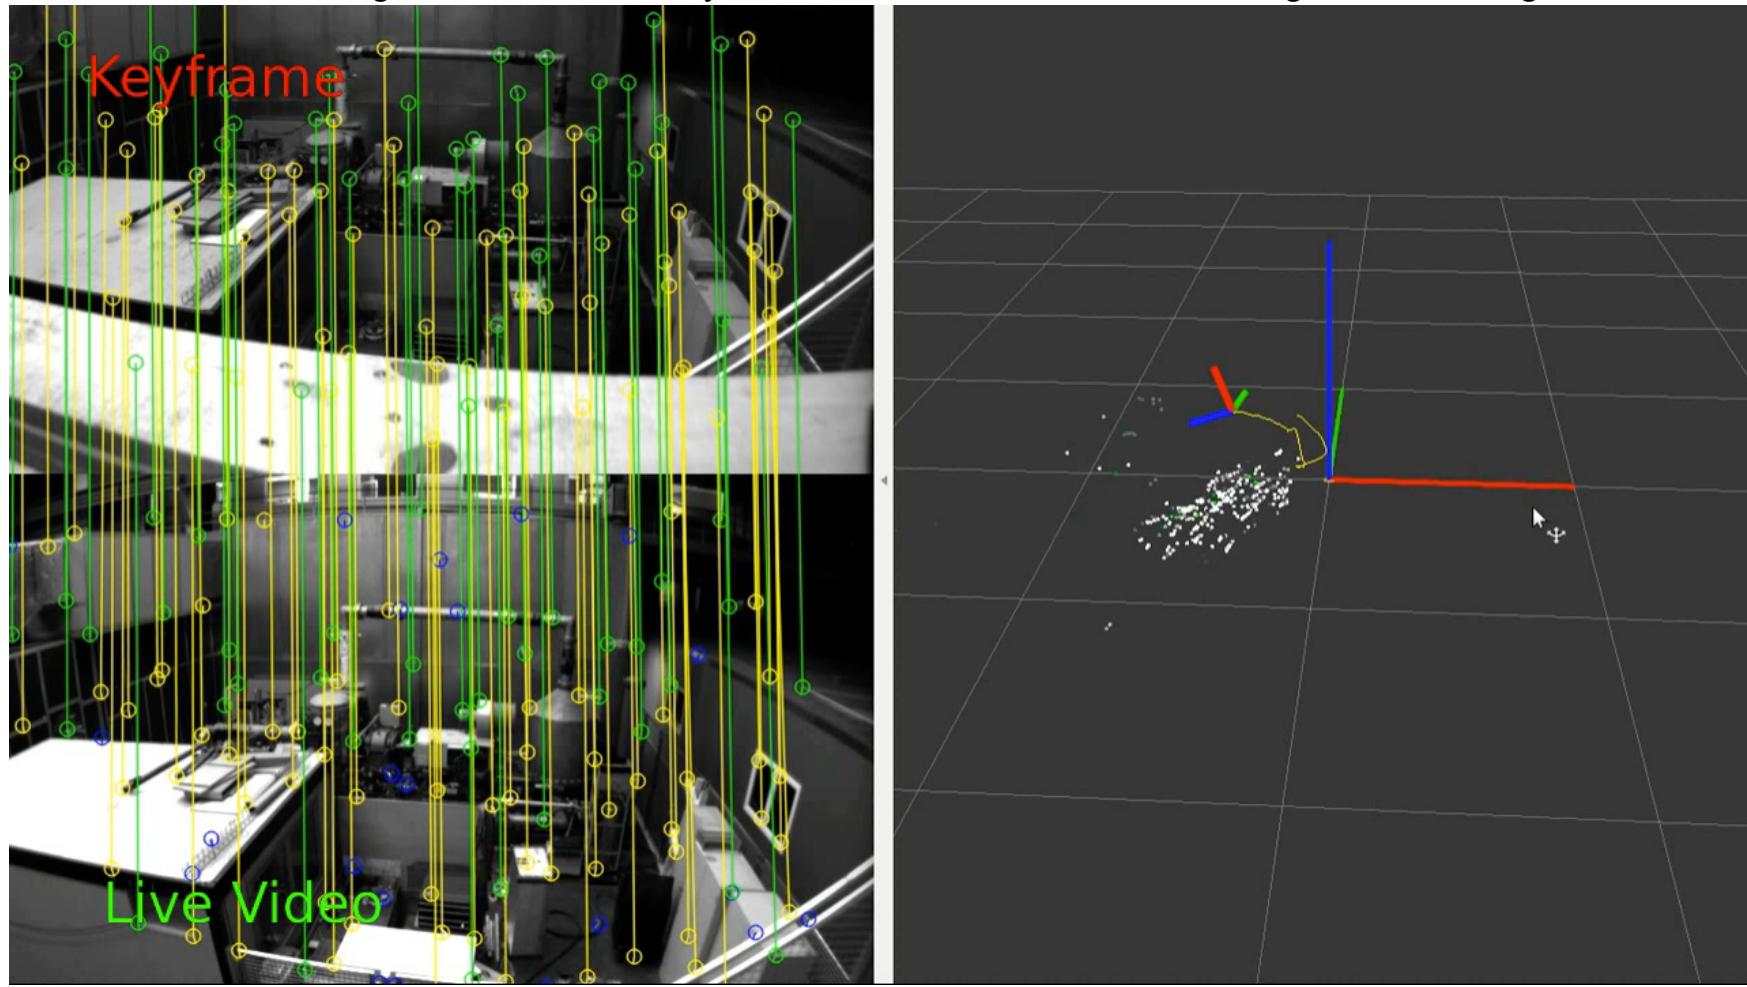
Today

1. Keypoint Detection, Description and Matching
2. Bootstrapping
3. Loop-Closure
4. Dense SLAM



OKVIS: Open Keyframe-based Visual-Inertial SLAM

Joint work with P. Furgale, M. Bosse, S. Lynen, M. Chli, V. Rabaud, K. Konolige, and R. Siegwart

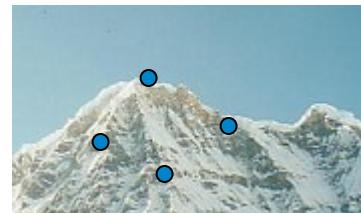


https://www.youtube.com/watch?v=TbKEPA2_m4

Keypoints: What Do We Want to Achieve?

Challenge 1:

Detector repeatability



bad

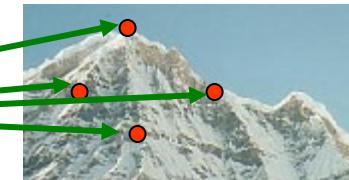
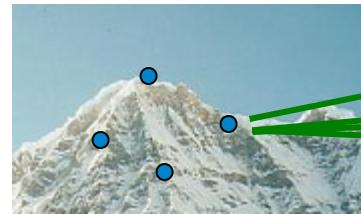


Challenge 2:

Descriptor quality

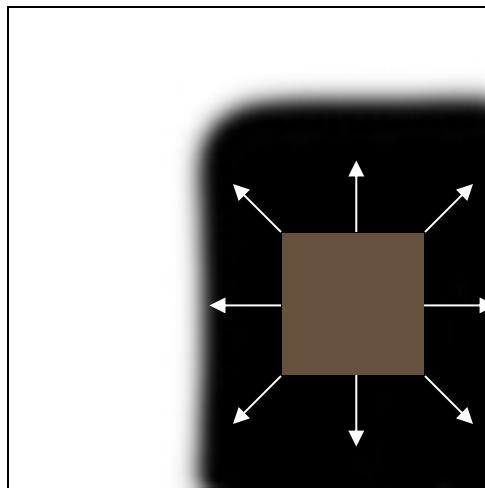
More Challenges:

- Speed: Detection and description / Matching, database retrieval
- Invariance with respect to rotation and/or scale?
- Tolerance to brightness / contrast changes
- Distribution in the image

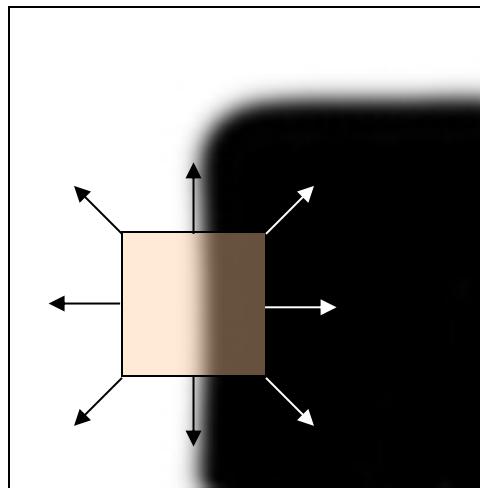


Identifying Corners

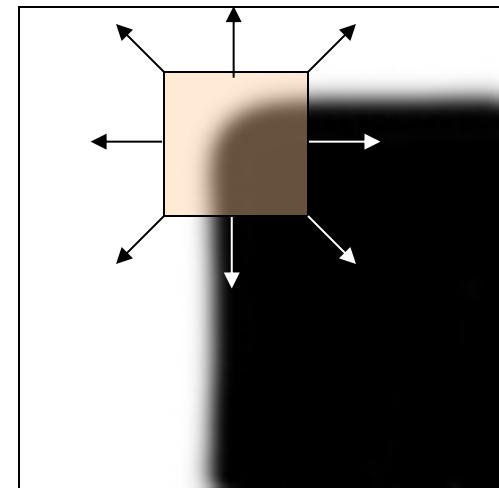
Shifting a window in **any direction** should give a **large change** in intensity in at least 2 directions



“flat” region:
no intensity
change



“edge”.
no change along the
edge direction



“corner”.
significant change in
at least 2 directions

How Do We Implement This?

$$SSD(\Delta x, \Delta y) \approx \sum_{x,y \in P} (I_x(x, y)\Delta x + I_y(x, y)\Delta y)^2 = [\Delta x \quad \Delta y] M \begin{bmatrix} \Delta x \\ \Delta y \end{bmatrix}$$

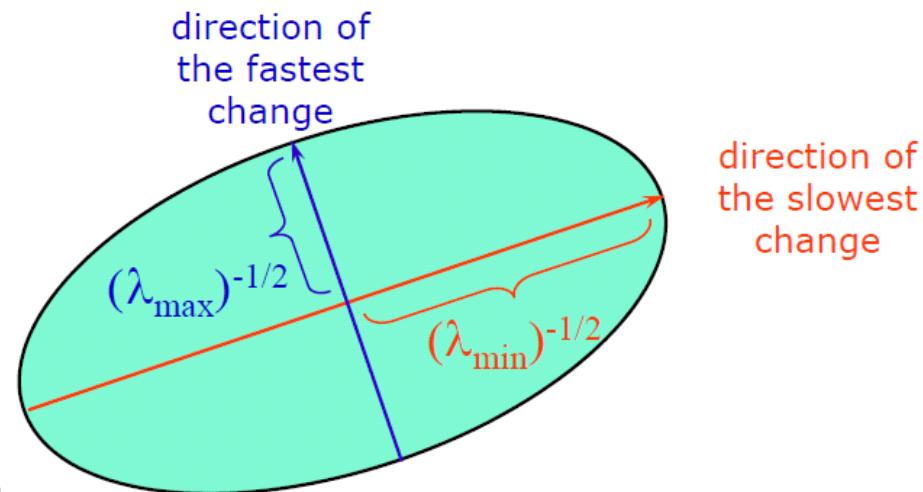
Intensity gradient along y-direction
Intensity gradient along x-direction

M is the “Second Moment Matrix”

$$M = \sum_{x,y \in P} \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix}$$

Since M is symmetric, we can rewrite M:

$$M = R^T \begin{bmatrix} \lambda_1 & 0 \\ 0 & \lambda_2 \end{bmatrix} R$$



where λ_1 and λ_2 are the Eigenvalues of M.

Corner Response Function [Harris and Stephens, AVC88]

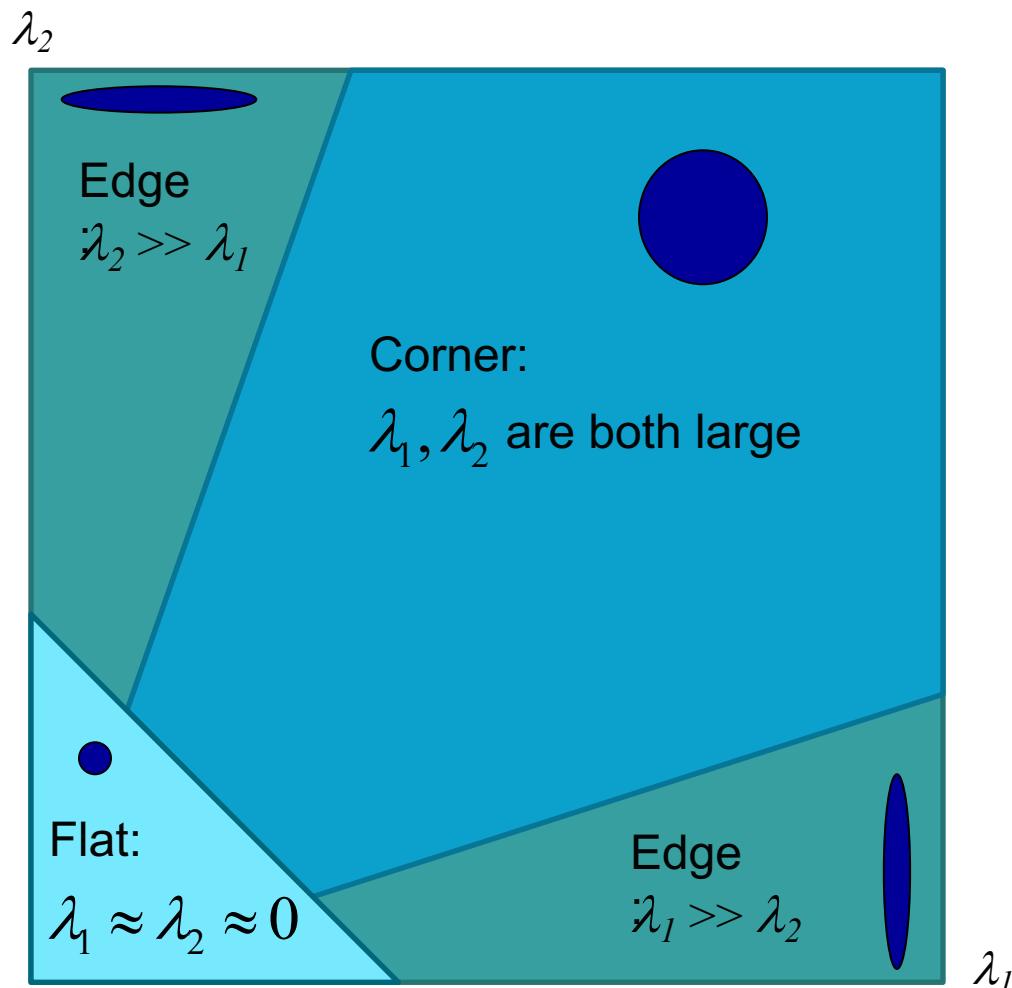
«Cornerness Function»

$$R = \det(M) - \kappa \operatorname{trace}(M)^2$$

$$= \lambda_1 \lambda_2 - \kappa(\lambda_1 + \lambda_2)^2$$

Where κ is a constant between 0.04 and 0.15.

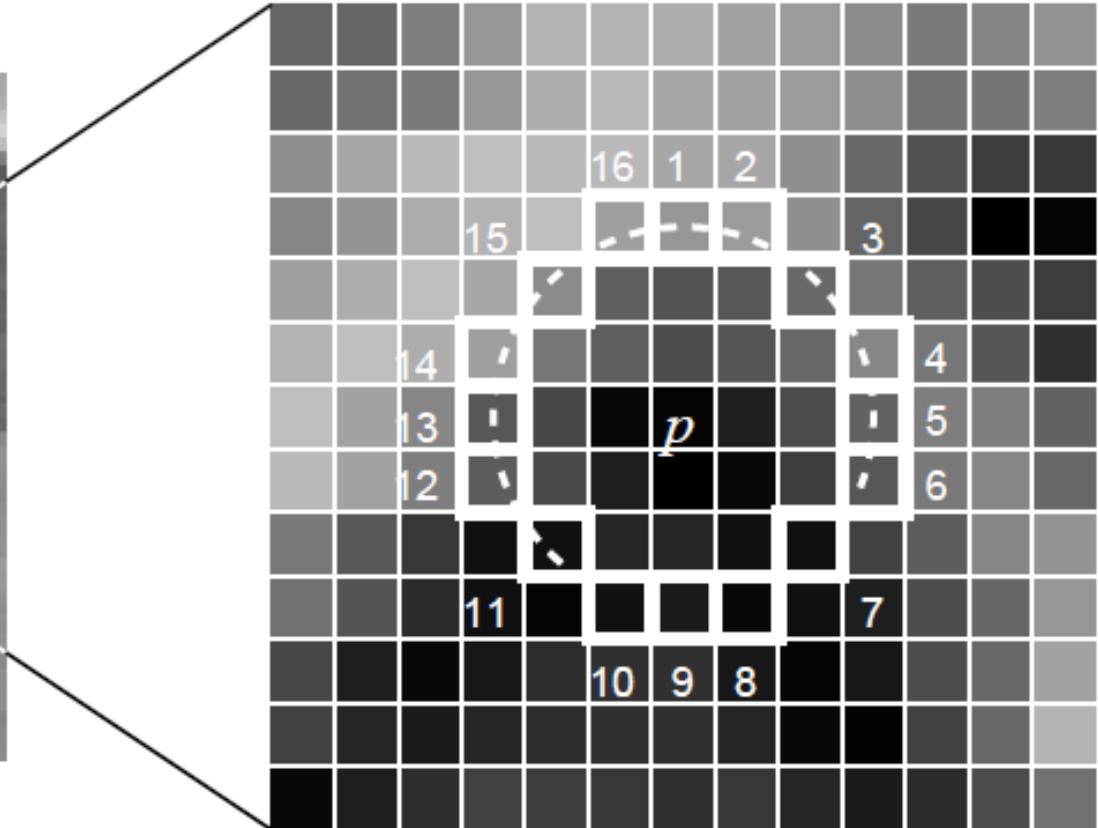
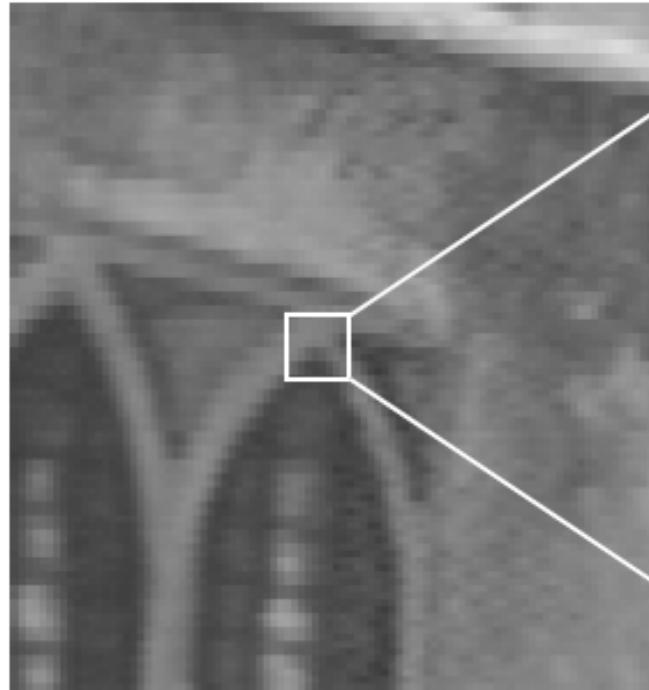
Note: this is invariant to rotation by definition.



FAST (-ER) Corners [Rosten et. al., PAMI2010]

Test: All points in circular segment brighter/darker by T than center (p)

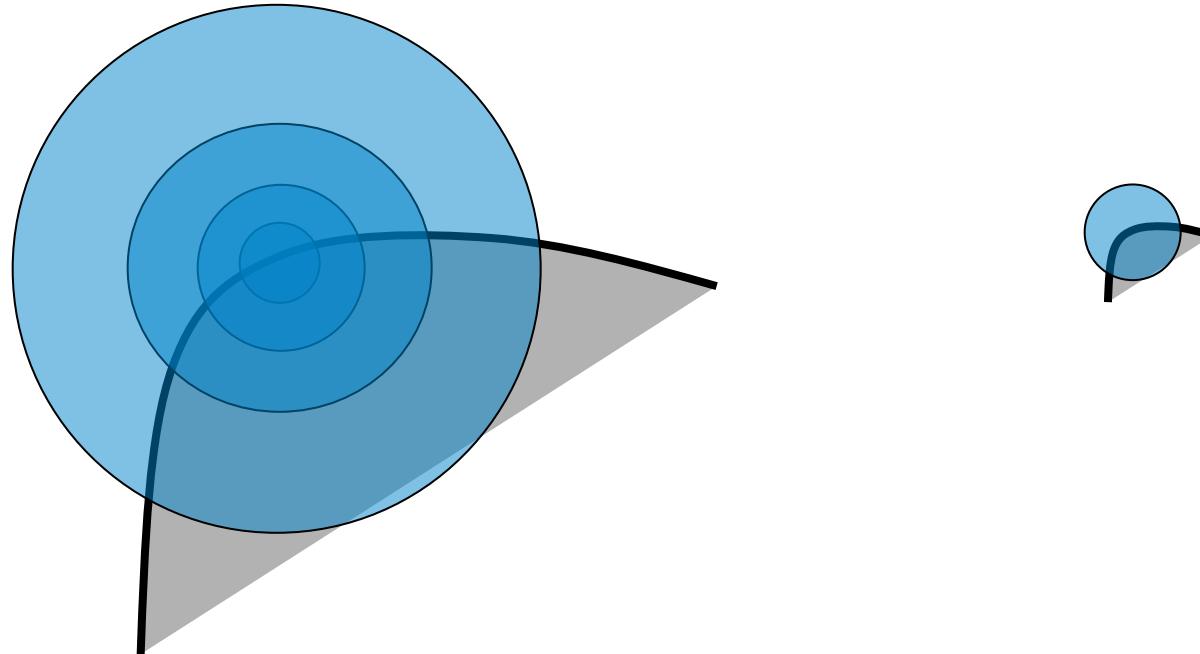
Very fast: in the order of 100 Mega-pixel/second



Scale Invariant Detection

Consider regions (e.g. circles) of different sizes around a point.

Corresponding regions will look the same in image space,
when the appropriate scale-change is applied



Excursion: Spatial Filters

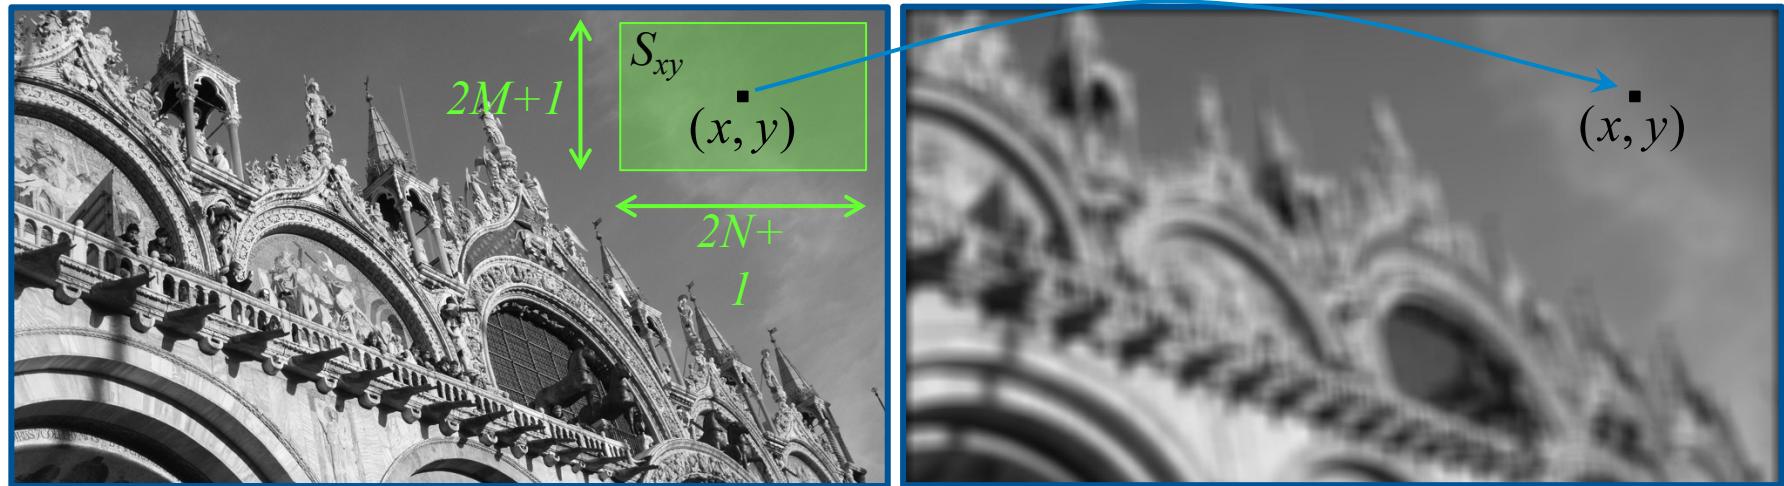


Image I

$$J(x, y) = \frac{\sum_{(r,c) \in S_{xy}} I(r, c)}{(2M + 1)(2N + 1)}$$

Filtered Image $J = F(I)$

Equivalent to **Convolution** of (windowed) **Kernel** k with the image:

$$J(x, y) = k(x, y) * I(x, y)$$

Blob Detection: LoG and DoG

Scale(t)-normalized Laplacian of Gaussian (LoG):

- Image smoothed by a Gaussian kernel:

$$L = g(x, y, t) * I(x, y)$$

- Apply Laplacian operator:

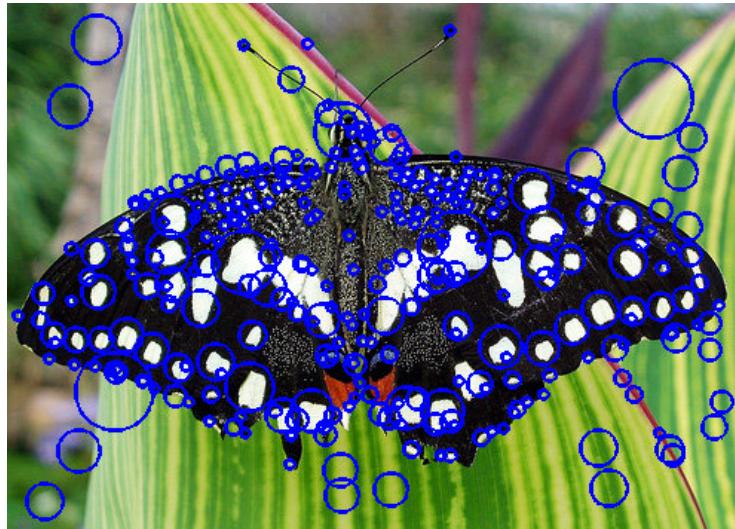
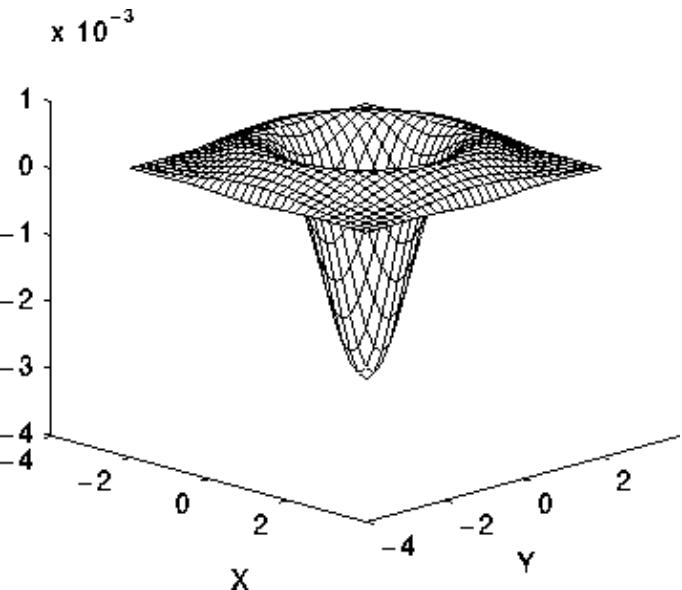
$$\nabla_{\text{norm}}^2 L = t \left(\frac{\partial^2 L}{\partial x^2} + \frac{\partial^2 L}{\partial y^2} \right)$$

Approximation:

Difference of Gaussians (DoG)

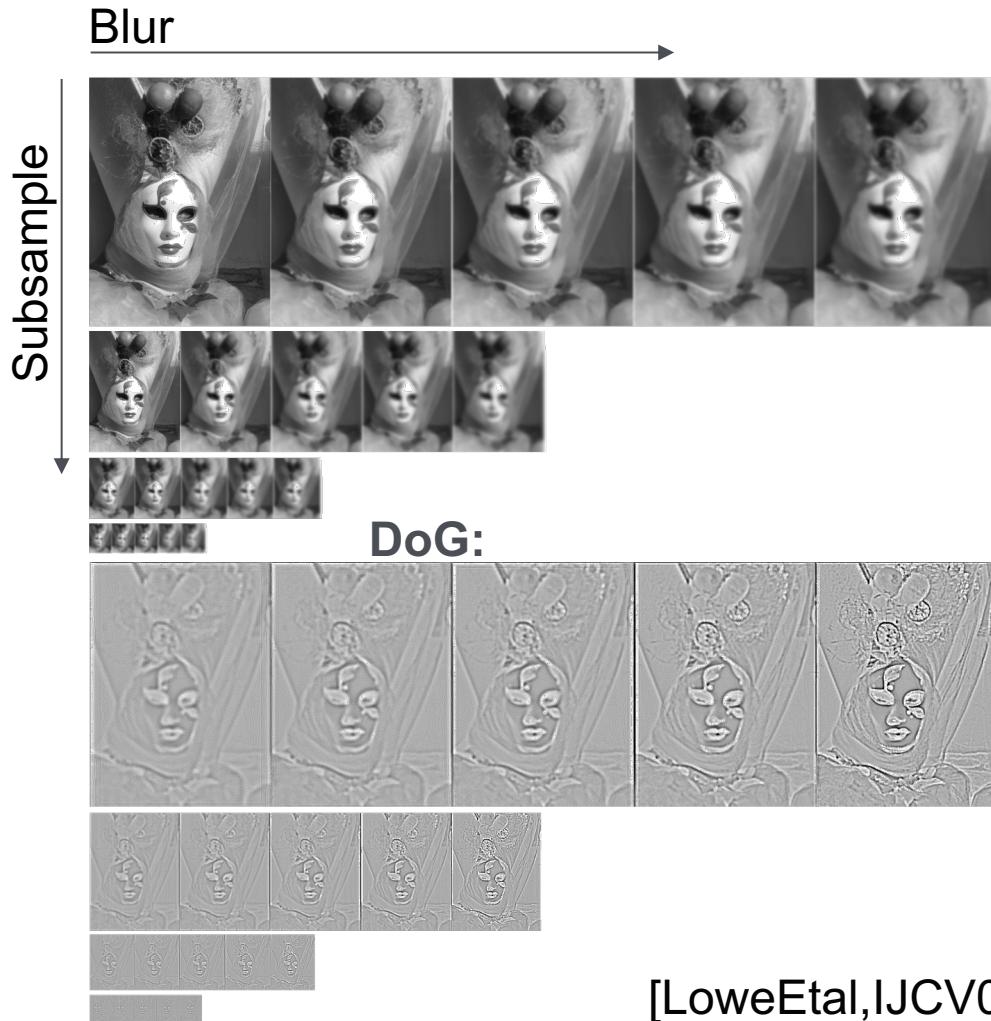
SIFT detector [Lowe et.al., IJCV04]

$$\Delta L = L(x, y, t) - L(x, y, kt)$$

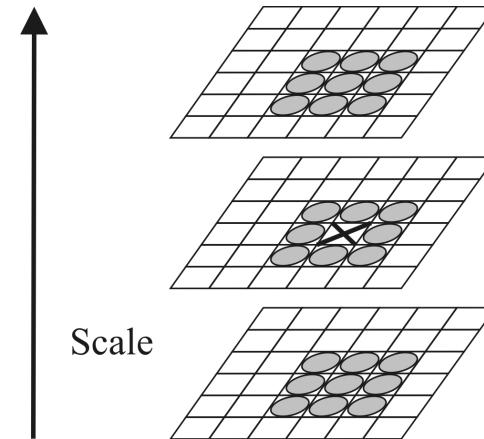


Micah Taylor, www.kixor.net, accessed June 2012

SIFT Detector (Keypoint Location + Scale)



1. Scale-space pyramid: subsample and blur original image
2. Difference of Gaussians (DoG) pyramid: subtract successive smoothed images
3. Keypoints: local extrema in the DoG pyramid



[LoweEtal,IJCV04]

SIFT: Orientation Assignment

1. Find ‘orientation’ of keypoint to achieve **rotation invariance**
2. Sample intensities around the keypoint
3. Compute a histogram of orientations of intensity gradients
4. Peaks in histogram: dominant orientations
5. Keypoint orientation = histogram peak
6. If there are multiple candidate peaks, construct a different keypoint for each such orientation

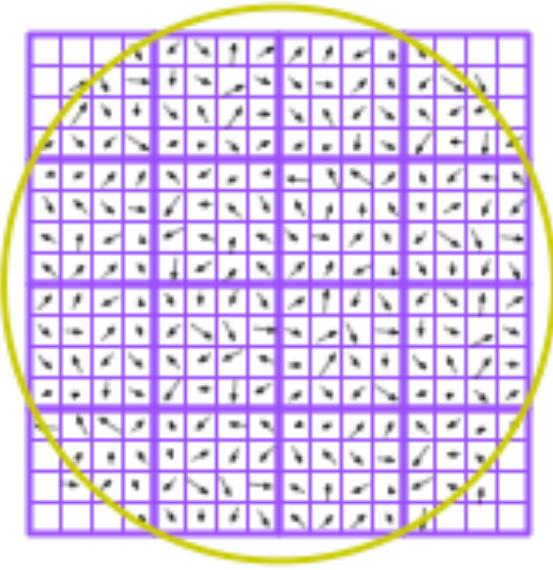
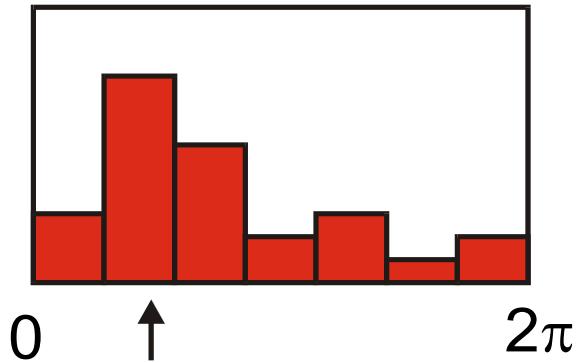
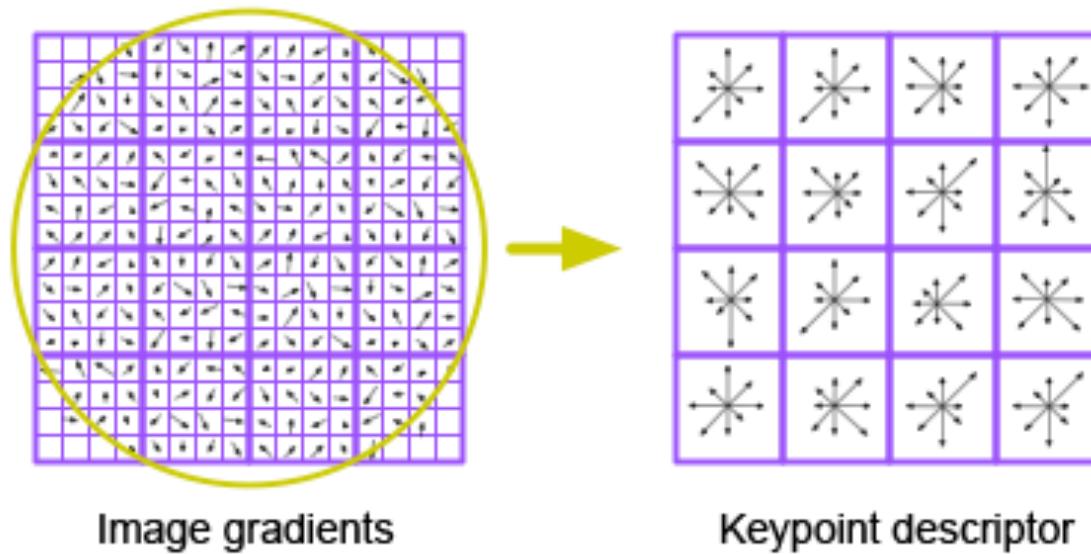


Image gradients



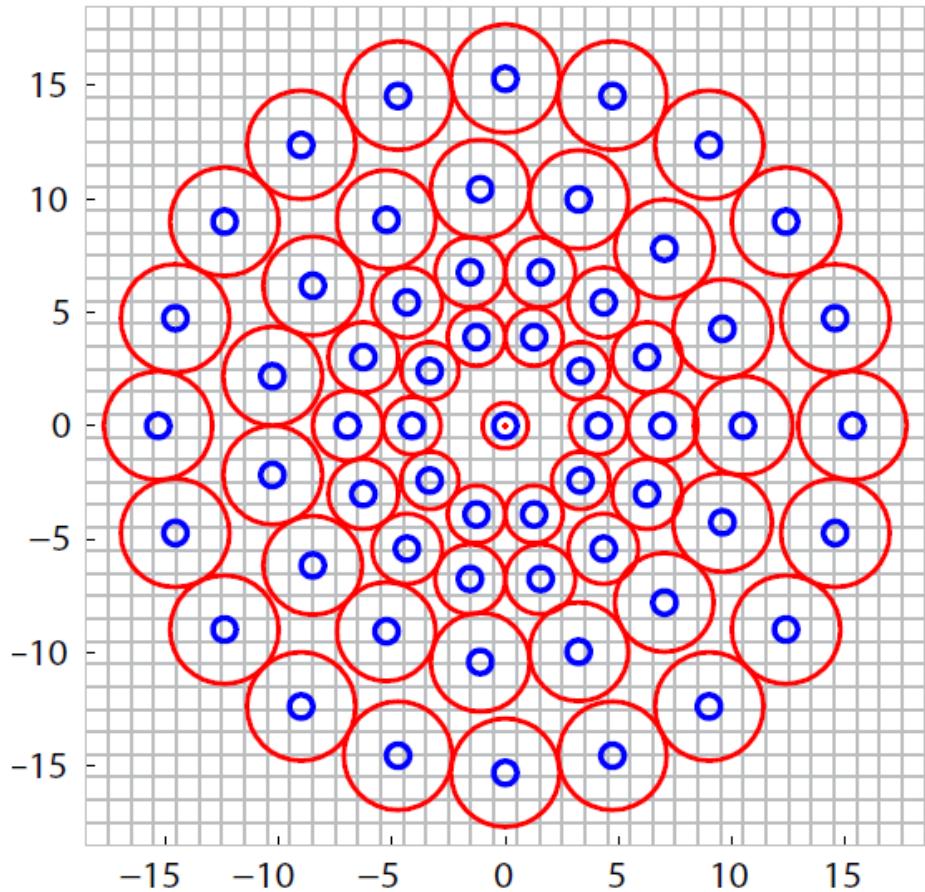
SIFT Descriptor

- We want it to be tolerant/invariant to certain changes, but yet very discriminative
- Describe gradient orientations relative to the keypoint orientation
- Divide keypoint neighborhood into 4×4 regions and compute orientation histograms along 8 directions
- SIFT descriptor: concatenation of all $4 \times 4 \times 8$ ($=128$) values
- **Matching:** L_2 -distance between these descriptor vectors



BRISK Descriptor [Leutenegger et. al ICCV 2011]

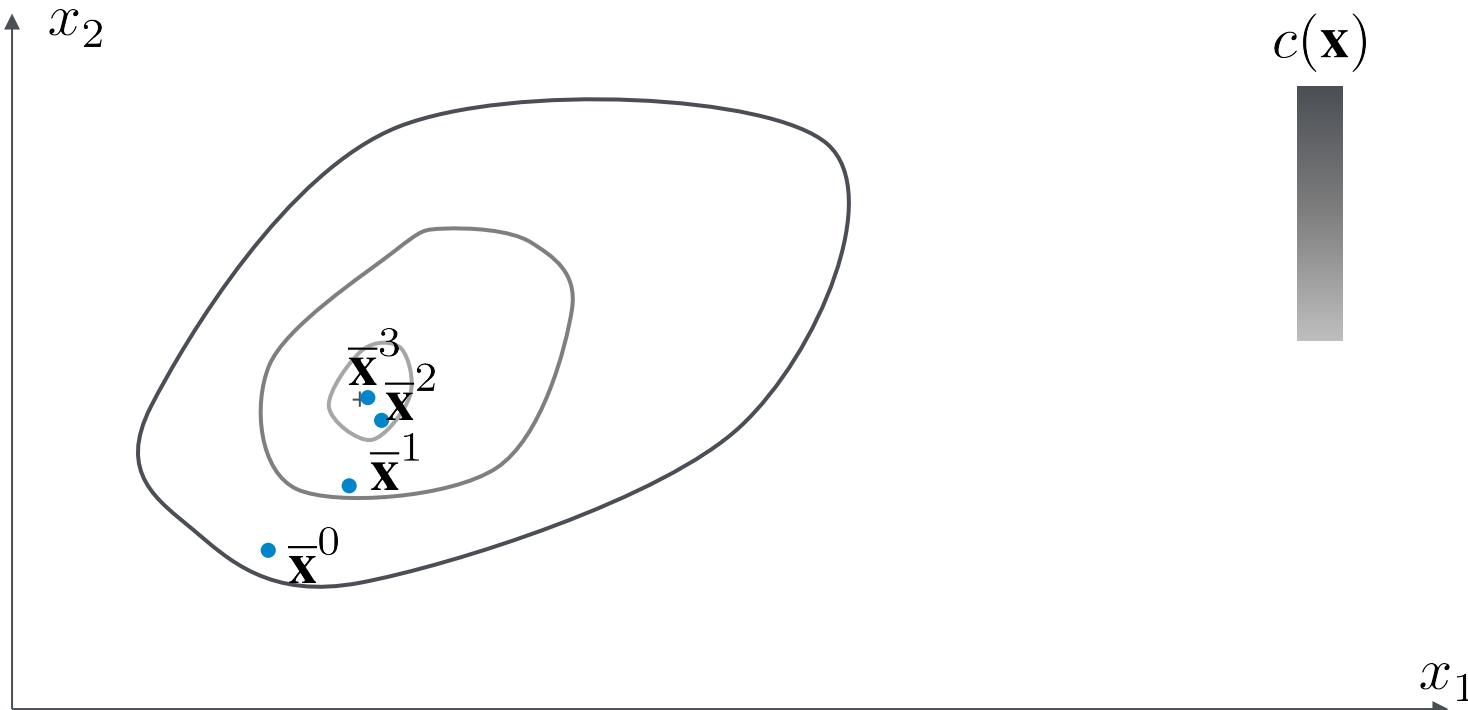
- Pattern used to access intensities in a keypoint neighborhood
- Red circles indicate the size of the smoothing kernel applied.
- Scaled and rotated versions stored in a look-up table
- Sampling point locations and smoothing is configurable
- Used for orientation assignment and brightness comparisons for **binary descriptor**
- Detection and descriptor speed: ≈ 10 times faster than SURF
- Ultra-fast **matching** using Hamming distance
(xor descriptors and count ones)



BRISK sampling pattern

How To Solve Nonlinear Least Squares Problems

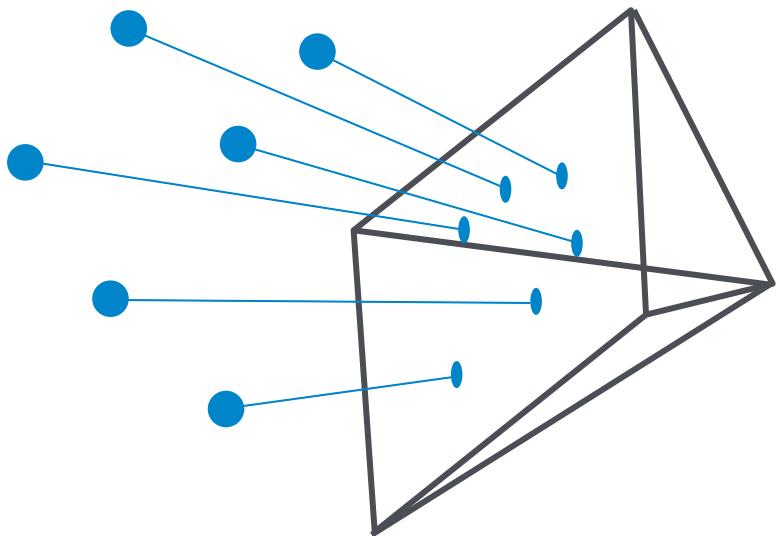
Iterative cost minimisation starting at an initial guess for \mathbf{x} :



But how do we find this initial guess?

3D-2D: The PnP Problem

Perspective n-Point Problem (also: Absolute Pose Problem): n 3D points are known and their 2D projections are observed in an image (known intrinsics)



Find: Solution for the full camera pose **directly** (i.e. no iterative solving).

Solutions: different approaches for various values for n ...

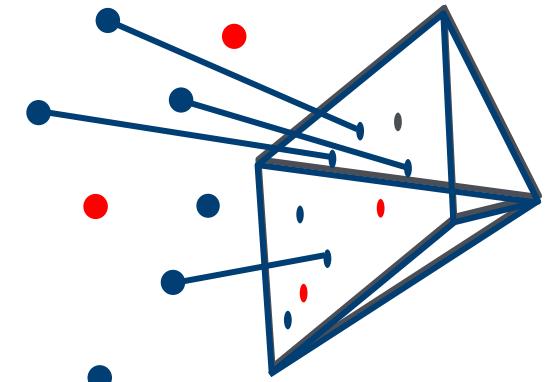
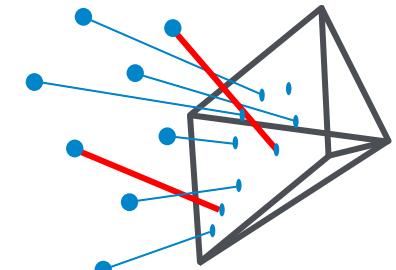
- State-of-the-art examples:
- UPnP (includes multi-camera setups)
[L. Kneip, H. Li, Y. Seo, ECCV'14]
- EPnP [V. Lepetit, F. Moreno-Noguer, P. Fua IJCV'08]

RANSAC: RANdom SAmpling Consensus

Problem: some amount of **wrong 3D-2D correspondences** called **outliers** may have a catastrophic effect PnP solution...

Find: Outliers and correct pose. **A Solution:** RANSAC.

```
function PnP(Ransac(points_3d, detections_2d, n, k, t, d)
besterr = inf
while iterations < k:
    maybeinliers <- select n points at random
    compute PnP solution (maybemodel) using maybeinliers
    for each point not in S:
        if point projection fits maybemodel with error < t:
            add point to alsoinliers
    if no. elements in alsoinliers > d
        inliers <- maybeinliers and alsoinliers
        compute bettermodel from inliers
        err <- mean reprojection error
        if err < besterr:
            bestfit = bettermodel
            besterr = thiserr
return [bestfit, inliers]
```



Stereo Triangulation

How do we initialise 3D points now that we know our camera pose?

Given: two corresponding 2D - points in 2 views (known poses), i.e. **two rays**.
 The problem is that in practice, 3D lines won't intersect exactly...

Find: a good hypothesis for the respective 3D point **directly**.

A (fast) solution: Midpoint Method

1. Find the points along the ray with minimum distance to each-other:

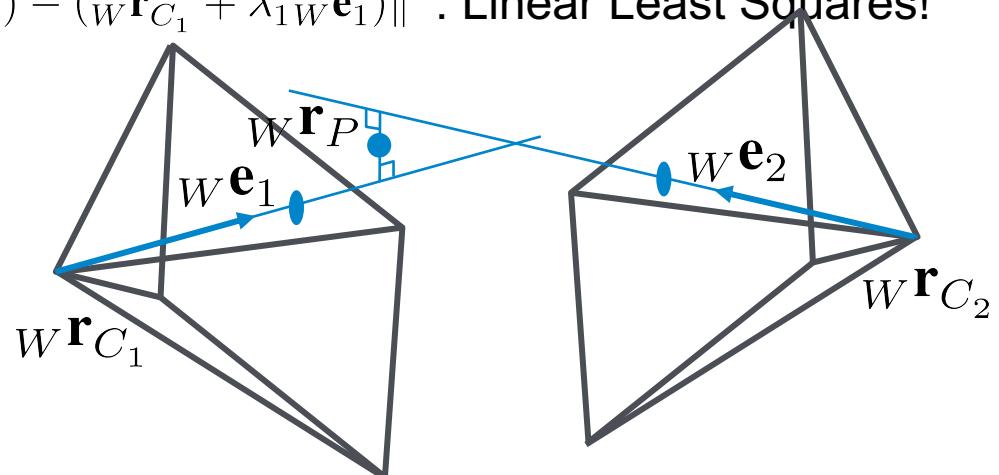
$$\lambda = \begin{bmatrix} \lambda_1 \\ \lambda_2 \end{bmatrix} = \operatorname{argmin} \|({}_W\mathbf{r}_{C_2} + \lambda_2 {}_W\mathbf{e}_2) - ({}_W\mathbf{r}_{C_1} + \lambda_1 {}_W\mathbf{e}_1)\|^2 : \text{Linear Least Squares!}$$

Solve the Normal Equations:

$$\mathbf{A}\lambda = \mathbf{b},$$

$$\mathbf{A} = \begin{bmatrix} 1 & -{}_W\mathbf{e}_1^T {}_W\mathbf{e}_2 \\ -{}_W\mathbf{e}_1^T {}_W\mathbf{e}_2 & 1 \end{bmatrix},$$

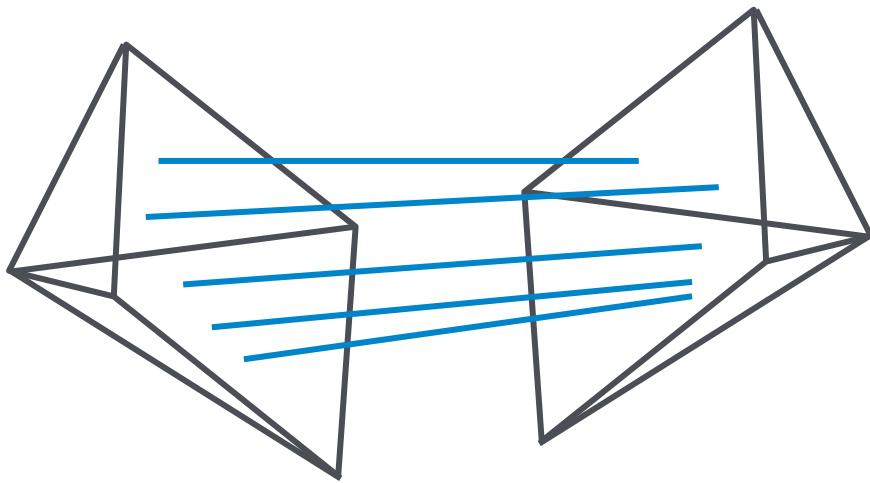
$$\mathbf{b} = \begin{bmatrix} \mathbf{e}_1^T ({}_W\mathbf{r}_{C_2} - {}_W\mathbf{r}_{C_1}) \\ -\mathbf{e}_2^T ({}_W\mathbf{r}_{C_2} - {}_W\mathbf{r}_{C_1}) \end{bmatrix}.$$



2. Pick the mid-point ${}_W\mathbf{r}_P = \frac{1}{2}(({}_W\mathbf{r}_{C_2} + \lambda_2 {}_W\mathbf{e}_2) + ({}_W\mathbf{r}_{C_1} + \lambda_1 {}_W\mathbf{e}_1))$.

2D-2D RANSAC

How about if we don't know any 3D points (e.g. monocular SLAM initialisation)?



A great library to that implements all those 3D-2D and 2D-2D algorithms (with RANSAC):

<http://laurentkneip.github.io/opengv/>

[9] L. Kneip, P. Furgale, "OpenGV: A unified and generalized approach to real-time calibrated geometric vision", Proc. of The IEEE International Conference on Robotics and Automation (ICRA), Hong Kong, China. May 2014

Various algorithms to be used with RANSAC:

- 5-point algorithm by Stewenius
- 5-point algorithm by Nister
- 5-point algorithm to solve for rotations directly by Kneip
- 6-point algorithm by Henrik Stewenius for generalised relative pose
- 7-point algorithm
- 8-point algorithm by Longuet-Higgins
- 17-point algorithm by Hongdong Li

Place Recognition

How should we find out that we see something previously seen?

A popular approach: **Bag of Visual Words**.

Training

1. Build a **Vocabulary** from training descriptors by clustering close descriptors into «words».

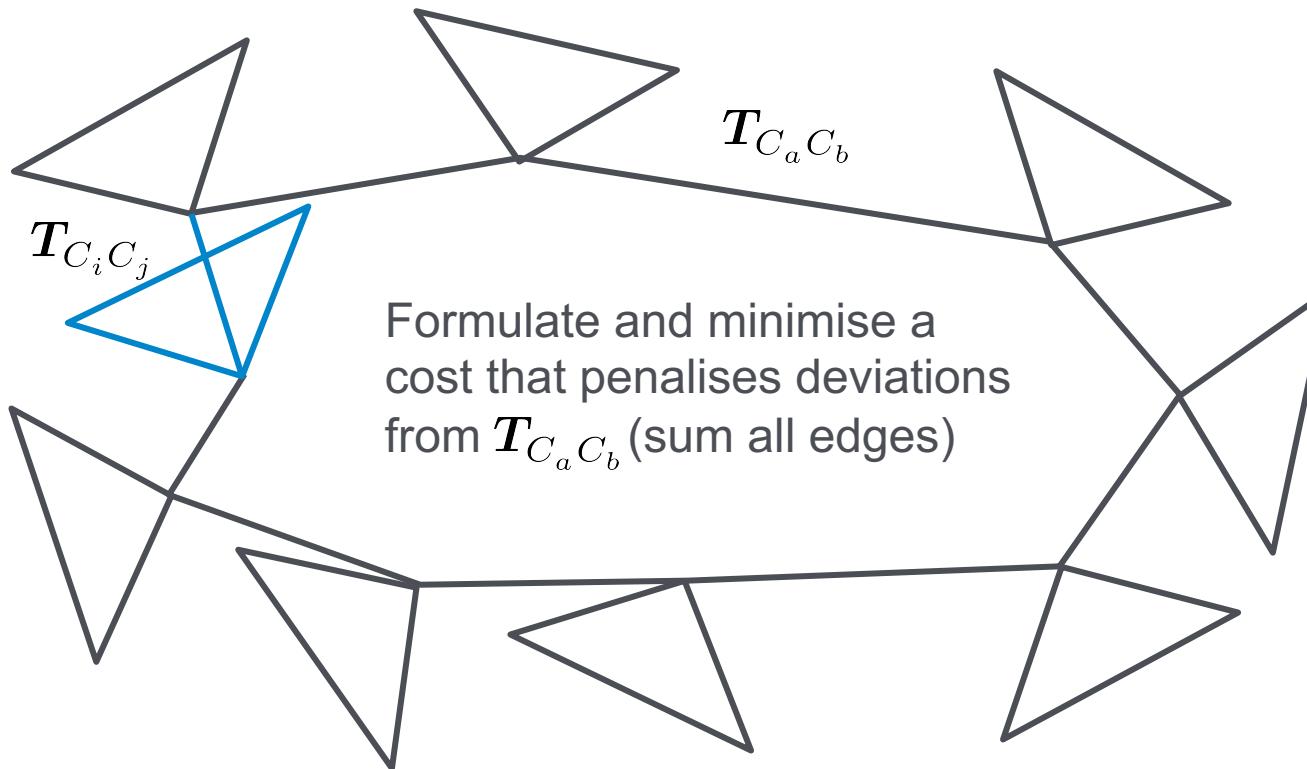
Runtime – for each keyframe:

1. Detect **keypoints** and extract **descriptors**.
 2. Build **histogram of visual words**, and perform some normalisation/re-weighting
 3. Database **query**: compute histogram similarity score for each location in the database.
 4. If it you find a match:
Apply individual **keypoints matching** and **verify** geometrically using some RANSAC which also gives you the **relative pose**.
Else insert the place into the database.
-

Pose Graph Optimisation

If trajectory / map needs to stay consistent after loop closure: bend it into place.

Example: Pose Graph Optimisation (Relaxation)



Dense SLAM Overview

Depth Cameras



Input: RGB image + Depth image

Principles to obtain depth image:

- Project pattern from baseline
- Passive stereo
- Time of Flight (ToF)

Typical SLAM algorithm:

1. **Track:** align RGB (and D) image with map to compute 6D pose.
2. **Map:** fuse new RGB-D into map.

Mono Cameras

Input: one image at the time only!

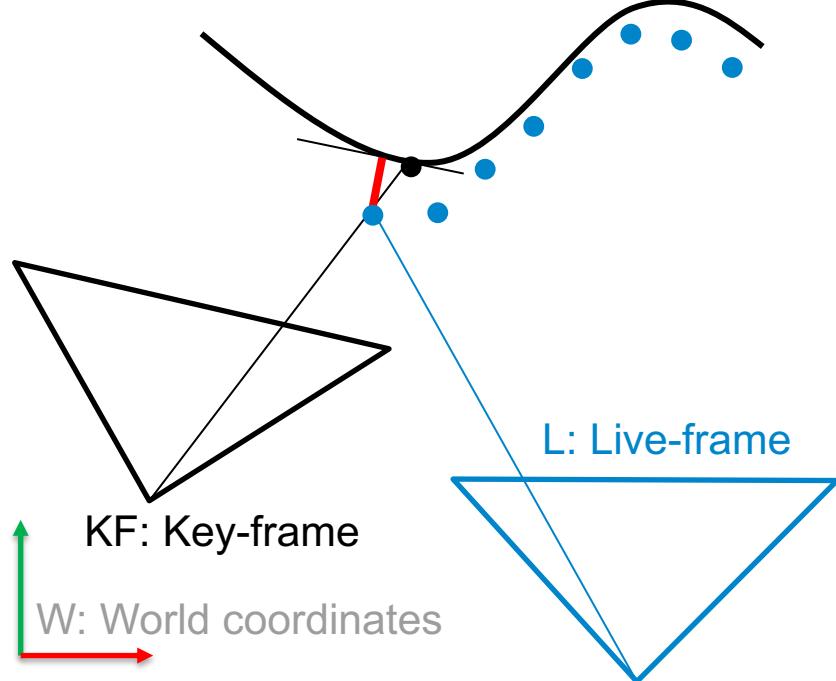
Typical SLAM algorithm:

1. **Track:** align image with map to compute 6D pose.
2. **Map:** fuse new data.
Problem: we don't know depth...
Options:
 - Compute depth maps from successive poses and then do the same as in RGB-D SLAM
 - Optimise the existing map directly to match the image(s)

Dense SLAM: Tracking



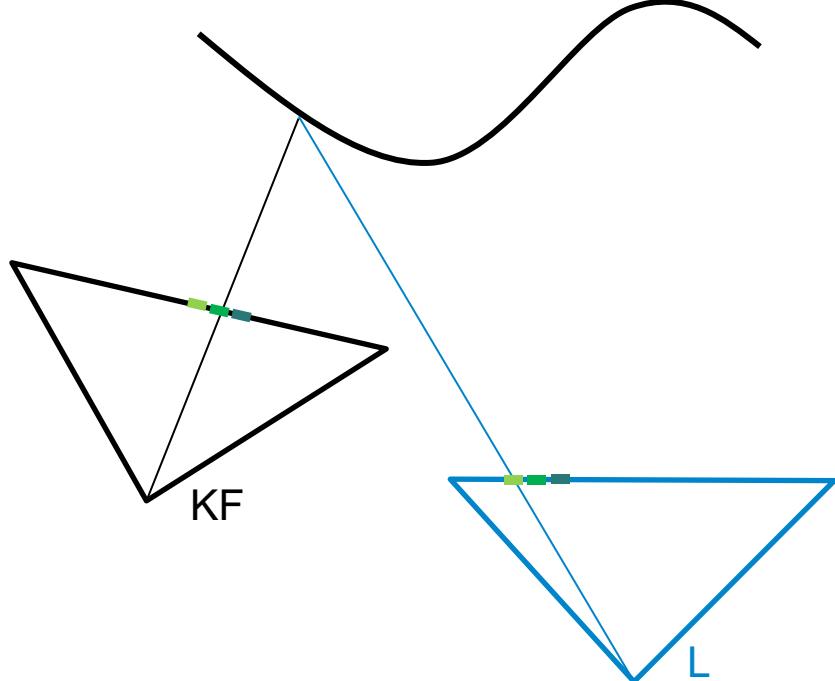
ICP error (geometric distance):



$$\mathbf{u}_{KF} = \pi(\mathbf{T}_{WC_{KF}}^{-1} \mathbf{T}_{WC_L} (\pi^{-1}(\mathbf{u}_L, D_L[\mathbf{u}_L])))$$

$$e_{ICP} = {}_W\mathbf{n}_{KF}[\mathbf{u}_{KF}] \cdot (\mathbf{T}_{WC_L} \mathbf{v}_L[\mathbf{u}_L] - {}_W\mathbf{v}_{KF}[\mathbf{u}_{KF}]) \quad e_p = I_{KF}[\mathbf{u}_{KF}] - I_L[\mathbf{u}_L]$$

Photometric error (colour difference):

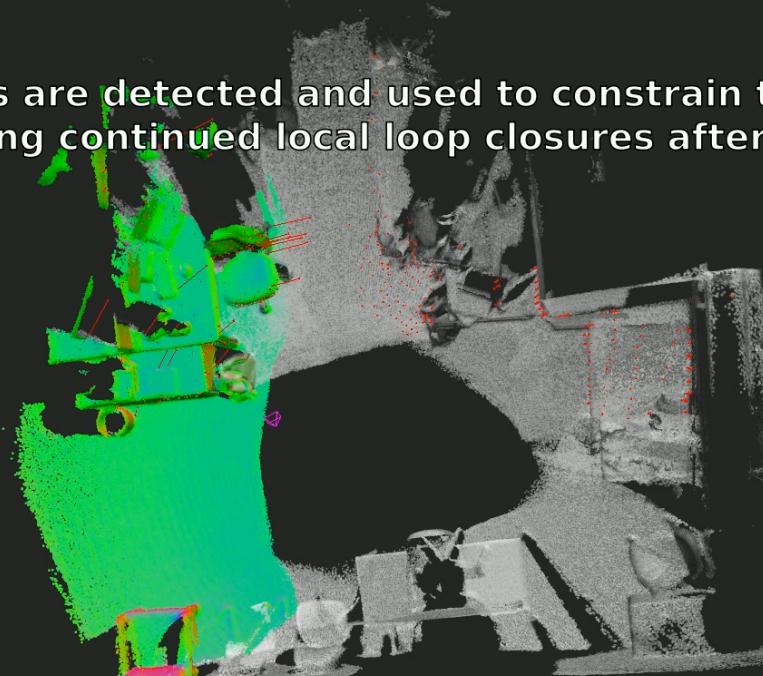


$$\mathbf{u}_L = \pi(\mathbf{T}_{WC_L}^{-1} \mathbf{T}_{WC_{KF}} (\pi^{-1}(\mathbf{u}_{KF}, D_{KF}[\mathbf{u}_{KF}])))$$

ElasticFusion (Again)

Overview (Real-time)

Global loop closures are detected and used to constrain the surface globally enabling continued local loop closures afterwards



Live Colour



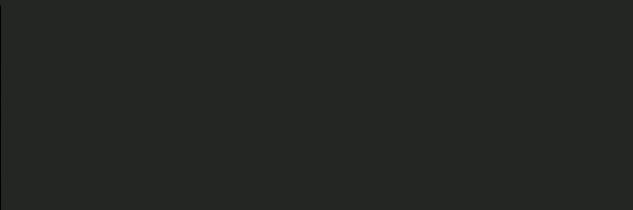
Live Depth



Active Colour

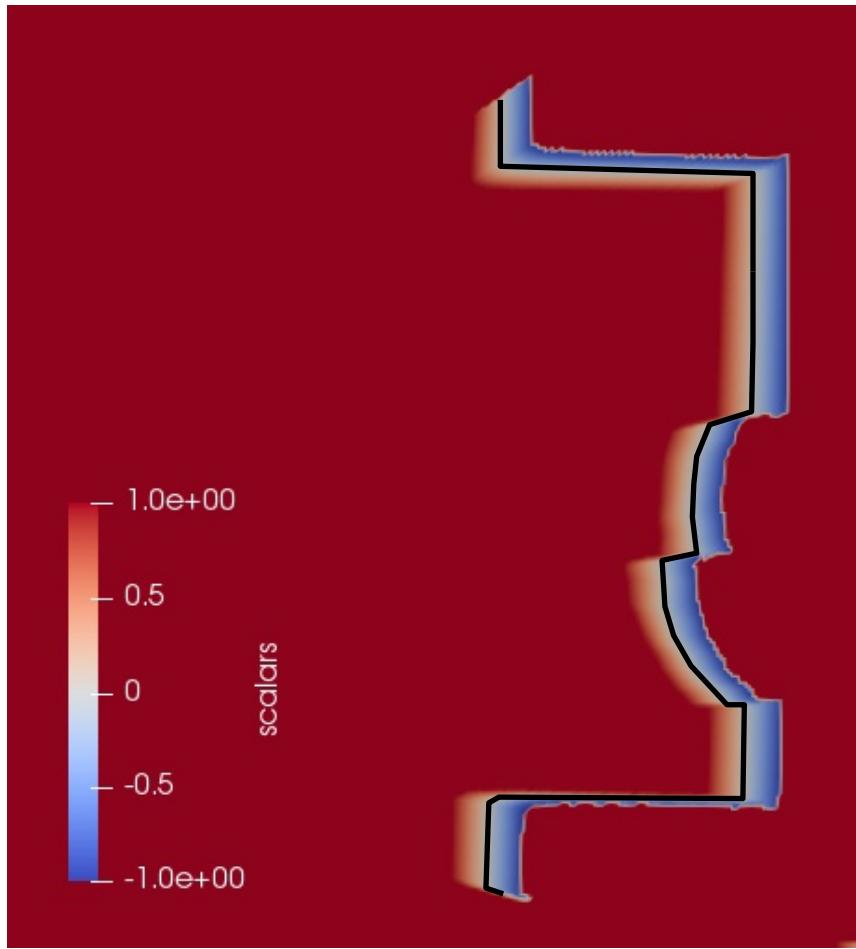


Active Surface



Dense Volumetric Mapping

- Space represented as **truncated signed distance field (TSDF)** or Log-Odds **occupancy values**
- **Surfaces** represented implicitly as **zero-crossings**
- Volume may be encoded in an **octree** to exploit environment's sparsity.



Any Questions?

See you on Thursday at 1pm for the practical.